



Mind, Market, and Mandate: A Field-Theoretic Synthesis for PPE

[T]-Theory Volume: Philosophy, Politics, and Economics

Alistair Johnson

Independent Researcher, Zurich, Switzerland

ORCID: 0009-0007-2194-0850

[T]-Theory: PPE

P14 · Capstone · For: The PPE Scholar

The Mandate Consensus Propagator — democratic field dynamics.

Politics as field theory:

- Electorate = 8D BRECVEMA field over social space
- Mandate = dominant attractor basin in the collective field
- Electoral volatility = field temperature T_{field}
- Consensus = $k_{\text{pol}} \rightarrow 0$ (long-range correlation)

Philosophy + Economics:

- Justice = rights invariant from Book 13 (Law)
- Welfare = negative Hopfield energy $-H(\Psi)$
- Political will = volitional injection $J_{\text{user}}(t)$ at scale 9
- Policy = reshaping the collective attractor landscape

I · The Mandate Consensus Propagator

The collective decision field in a democratic system:

$$\Psi_{\text{pol}}(\mathbf{x}, t) = \sum_i w_i \Psi_i(t) \delta(\mathbf{x} - \mathbf{x}_i)$$

The **mandate** at time t = the dominant eigenmode of the collective field Ψ_{pol} . Consensus = when one eigenmode captures $> 50\%$ of field energy.

The propagator of democratic consensus:

$$G_{\text{mandate}}(\tau) = e^{-\tau/\tau_{\text{mandate}}} \cos(\omega_{\text{election}} \tau)$$

τ_{mandate} = political mandate lifetime (days to years);
 ω_{election} = electoral cycle frequency.

II · Policy as Field Perturbation

A policy intervention P perturbs the collective field:

$$\delta\Psi_{\text{pol}}(t) = G_{\text{mandate}} * J_P(t)$$

Short-term policy impact = the field response over $\tau \ll \tau_{\text{mandate}}$. Long-term institutional change = basin reshaping (structural reform).

The political philosopher's problem: design $J_P(t)$ to move the collective field toward a desired attractor, with minimal coercion.

III · The Three Disciplines as One Field

- **Philosophy**: what is the optimal attractor? (normative H)
- **Politics**: how to reshape the basin? (dynamics of ∇H)
- **Economics**: what is the cost of each trajectory? (action S)

PPE = field theory with a normative Lagrangian. The USF provides the common mathematical language.

IV · Deliberative Democracy as Field Resonance

Deliberative democracy (Habermas, Rawls) $\equiv$ field resonance: citizens align their BRECVEMA fields through dialogue. Legitimacy = the degree of field coherence in the collective.

G-ID: *The Mandate Consensus Propagator* — collective decision kernel in democratic field theory. ORCID: 0009-0007-2194-0850
· CC BY 4.0 · Zurich 2026

Contents

The Green Propagator	10
Introduction: The Same Field Governs Mind, Market, and Mandate	10
Philosophy: Mind and the Field	11
Politics: Governance as Field Architecture	11
Economics: The Field as Market	12
The Patient Perspective	12
What This Book Offers the PPE Student or Scholar	13
The Programme	14
The Gap the Programme Addresses	14
The Structure of the Argument	15
The Method: Mathematical Co-identification	16
What It Is	16
Why It Matters as Method	16
When MCI Is Over: The Verification Threshold	17
The Model: The Soma-Field	19
Five Co-identifications	19
What the Model Predicts	20
The Empirical Test: QUANT-EXP-1	21
The Prediction	21
The Experiment	21
Results	21
The Penrose Connection	22
The Lived Case: Field Notes from the Inside	22
Extensions: Music, Film, and the Domain Generality of the Model	23
Music-Induced Affect	23
The Tensor: An Abstract Film	24
The Argument as a Whole	25
What Remains	26
Data and Code Availability	27

Introduction	28
Nash Equilibrium = Hopfield Minimum	28
Coordination Games and the Multi-Attractor Regime	29
Market Crashes as Phase Transitions	29
The FM-HN Extension: Regulatory Intervention	30
The Prisoner's Dilemma as Topological Obstruction	30
Applications	30
Conclusion	31
Introduction	32
Laws as Attractor Basin Constraints	32
Rights as Topological Invariants	32
Regulation as $J_{law}(t)$	33
The Rule of Law as Ergodicity	33
Legal Uncertainty and Attractor Fragmentation	34
Constitutional Law as Meta-Constraints	34
Applications and Predictions	35
Conclusion	35
A Note on Method	36
Introduction: The Inadequacy of Existing Maps	36
Background	37
Lived Experience as a Research Position	37
The Body-Mind Problem in Clinical Practice	38
The Felt Sense and Sub-Perceptual Emotion	38
Quantum Field Theory: Structure, Not Metaphor	38
Hopfield Networks and the Energy Function	39
The Soma-Field Model	39
Emotions as a Persistent Wave Field	39

The Perception Threshold	40
A note on the intelligence quotients	41
A note on brane thickness	41
The Interaction of Emotional Modes	42
The Energy Landscape	42
The Hopfield Energy Function	42
Attractor States: Fight, Flight, Freeze, and Regulated Calm	42
Dissonance and Resolution	44
The Neurodivergent Field: ASD, ADHD, and C-PTSD as Operator Modifications	44
Complex PTSD: Memory Kernel and Asymmetric Coupling	44
ADHD: High Temperature, Low Damping, Pink Noise	45
Autism Spectrum Condition: Sparse Coupling and Modified Projection	45
The Composed Condition	46
The Soma-Field Instrument	46
Rationale	46
Design	46
The Feedback Loop	47
Pluggable Emotion Models	47
Clinical Implications	47
Assessment	47
Intervention	48
Psychoeducation	48
Neurodivergent Profiles as Structural Realities	48
Limitations and Future Directions	48
Conclusion	49
Introduction: The Missing Unification	50
The Green's Function as the Universal SHO	50
The String Theory Problem	51
The Identification	51
Consequences	52
The 11-Dimensional Architecture	52
The Decomposition	52

The Limbic Axis as the Horava-Witten Orbifold	53
The 20-Scale Dial	53
The Organism Hierarchy	54
Three Tiers	54
Consciousness as Phase Transition	55
The Classical Gap	55
The Threshold	55
What Consciousness Is	55
Relation to Existing Frameworks	56
McFadden’s CEMI Theory	56
Schreiber’s Modal Homotopy Type Theory	56
The Σ -Type Formulation of the USF	56
Hoffman’s Conscious Agents	57
Formal Verification	57
The Volitional Agent	58
From Autonomous to Driven Dynamics	58
The Somatic Injection	58
Patient to Pilot	59
Discussion	59
What Has Been Claimed	60
The Correspondence Principle at Every Scale	60
Lean 4 as Epistemological Standard	60
Conclusion	61
Introduction: The Two-Culture Problem Within Science	62
The Typeverse	62
The Procedure in Detail	63
Step One: Extract the Type Signature	63
Step Two: Search the Typeverse	64
Step Three: Identify and Import	64
The Formal Computational Structure: Abduction, Aesop, and the Loop	65
Historical Precedents	66
Veneziano (1968): The Bootstrap Amplitude and String Theory	66

Hopfield (1982): The Ising Hamiltonian and Neural Memory	66
Wilson (1971): Block Spins and the Renormalisation Group	67
Black and Scholes (1973): The Heat Equation and Options Pricing	67
Jaynes (1957): Thermodynamic Entropy and Bayesian Inference	67
Penrose (1971): Spin Networks and Spacetime Geometry	68
Selinger (2010): Linear Maps and Quantum Processes	68
The Soma-Field as a Worked Example	68
Co-identification I: The Conscious Percept as Green's Function	68
Co-identification II: The Attractor Landscape as Ising Hamiltonian	69
Co-identification III: The Perception Threshold as Brane Thickness	69
Co-identification IV: The Coupling Matrix as $G \square$ Manifold	70
Co-identification V: Therapeutic Processing as Renormalisation Group Flow	70
A Partial Map of the Typeverse	71
Propagator-Class Structures	71
Energy-Function-Class Structures	71
Topological-Class Structures	71
Renormalisation-Class Structures	72
Scattering-Class Structures	72
Einstein-Coefficient-Class Structures	72
Failure Modes	73
Type Coincidence Without Structural Identity	73
Non-Commutative Functors	73
Over-identification	73
The Metaphor Trap	74
Epistemological Status	74
What Co-identification Claims	74
Why It Is Not Analogy	74
The Role of Artificial Intelligence	75
The Methodology as Practice	75
Falsifiability Protocol for Publication Use	76
Minimal Registration Template	76
Disconfirmation Rules	76
Worked Registration Sketch (Soma-Field)	77
Reviewer-Facing Scope Labels	77
Negative Control: When Matching Units Is Not Enough	77

Worked External Example (Non-Soma): Black-Scholes to Heat Equation	78
Step A: Signature Match	78
Step B: Import Claim	78
Step C: Assumption Checklist	78
Step D: Imported Theorem and Prediction	79
Step E: Disconfirmation Condition	79
Replication Package Requirements	79
Reviewer-Risk Objections and Responses	80
Independent Replication Ledger Linkage	80
Conclusions	80
Acknowledgements	81
Introduction	82
Background	83
The Body-Mind Problem in Clinical Practice	83
The Felt Sense and Sub-Perceptual Emotion	84
Quantum Field Theory: Structure, Not Metaphor	84
Neural Network Energy Functions and Hopfield Networks	86
The Formal Correspondences: Where the Link Was Seen	88
The Body Schema, Interoception, and Pain	91
Correspondence with Existing Emotion Representations	92
The Soma-Field Model	94
Emotions as a Persistent Wave Field	94
The Perception Threshold	95
The Interaction of Emotional Modes	96
The Three-Layer Architecture	96
The Energy Landscape	99
The Structure of the Emotional Energy Function	99
Attractor States: Fight, Flight, Freeze, and Regulated Calm	99
The Coupling Matrix as a Personal Signature	100
Dissonance and Resolution	101
The Acoustic Analogy	101
The Resolution Principle	101
The Soma-Field Instrument	102

Rationale	102
Design	102
The Feedback Loop	103
The Pluggable Emotion Model	103
Clinical Implications	104
Assessment	104
Intervention	104
Psychoeducation	104
Neurodivergent Conditions as Operator Modifications	105
Limitations and Future Directions	108
Conclusion	109
Publication Claim Registry	110
Claim-Evidence-Result Matrix	111
Replication Package Requirements	112
Reviewer-Risk Objections and Responses	112
Independent Replication Ledger Linkage	113
Acknowledgements	113
Conclusion: The Integrated Account	114
The Common Structure	114
The Patient as Political-Economic Subject	115
What a Field-Theoretic PPE Would Teach	115
Open Questions for PPE	115

The Green Propagator

G-ID: *The Mandate Consensus Propagator — collective decision kernel in democratic field theory*

The Mandate Consensus Propagator is the collective decision kernel — the mathematical description of how individual preferences aggregate into a shared political will, how long that will persists, and under what conditions it dissolves. In this book you will encounter three disciplines that have historically spoken past each other: philosophy (what is the good?), politics (who decides?), and economics (at what cost?). The propagator is the common language. Watch for how the same mathematical object — a Green's function with characteristic timescale τ_{mandate} — appears in Rawlsian deliberation, in Arrow's impossibility theorem, and in market mechanism design. The question of collective action is, at its mathematical core, a question about the topology of this propagator.

Introduction: The Same Field Governs Mind, Market, and Mandate

Oxford PPE — Philosophy, Politics, and Economics — is an unusual curriculum. It was assembled in 1920 with the explicit intention that the three disciplines illuminate each other: that philosophical rigour improves political reasoning, that economic analysis sharpens political choices, that political reality disciplines philosophical abstraction, that philosophical foundations constrain economic modelling. In practice, the three disciplines have increasingly gone their separate ways. Philosophy of mind has little contact with macroeconomics. Political theory debates justice in near-total isolation from game theory. Economists build equilibrium models without engaging the philosophical foundations of rational choice. The integrative promise of PPE has been partially fulfilled in applied ethics and political economy, but the deepest integration — a single framework that governs all three domains from first principles — has not been achieved.

This book proposes that the Universal Somatic Field (USF) provides exactly that framework. The same master equation that describes the dynamics of felt experience also describes Nash equilibrium selection in strategic games, the stability of legal and political institutions, and the dynamics of market systems. This is not a claim that philosophy, politics, and economics are *reducible* to physics. It is the claim that they all instantiate the same *mathematical structure* — a structure that the USF framework makes explicit and whose implications can be derived, not merely asserted.

Philosophy: Mind and the Field

The philosophical pillar of PPE begins with questions of mind, agency, and value. What is a preference? What makes an action rational? What is the good life, and what grounds its value? These are questions that standard economics treats as primitives — preference orderings are given, rationality is defined axiomatically, utility functions are postulated — and that philosophy finds the economic treatment unsatisfying precisely because it is circular: it defines rationality as preference-satisfaction and preferences as what rational agents have.

The USF framework grounds preferences in field dynamics: a preference is a gradient in the somatic field energy landscape. What you prefer is what your somatic field moves toward; what you disprefer is what it moves away from. Rationality, in this framework, is not an axiom but a property: a rational agent is one whose field dynamics track the gradient of a consistent energy landscape. Inconsistent preferences (the preference reversals documented by Kahneman and Tversky, the intransitivities observed in experimental economics) are not anomalies to be explained away; they are signatures of a landscape that has been modified by context — different attractors are accessible in different framings.

This gives the philosophy of mind a connection to the economics of decision-making that neither has had before. The phenomenology of choosing — the felt quality of desire, aversion, deliberation, regret — is the inside view of a field navigating its energy landscape. The normative question (what should one prefer?) is the question of what attractor landscape is worth having — a question that the framework makes well-posed even if it does not answer it.

Politics: Governance as Field Architecture

The political pillar of PPE addresses the organisation of collective life: justice, legitimacy, authority, democracy, rights. The USF framework connects directly to the accounts developed in the companion law volume: political institutions are constraint fields that modify the somatic field energy landscape; rights are topological invariants; the rule of law is ergodicity; legitimacy is the condition under which the constraint field is accepted as a modification of the landscape rather than resisted as an external imposition.

The specifically political question that the framework illuminates is the **design problem**: given a desired collective outcome (a fair distribution, a stable democracy, a sustainable commons), what institutional architecture — what pattern of attractors, barriers, and channels in the social somatic field — will produce that outcome as a stable equilibrium? This is a field-design problem, and the USF framework provides the tools to address it: the WKB formula for minimum intervention strength, the topological stability analysis for rights, the ergodicity criterion for the rule of law.

The framework also addresses the pathological cases. Authoritarian capture — the concentration of power in the hands of a single actor who modifies the landscape in their own interest — corresponds to the non-ergodic case: the powerful actor's field is protected from accountability by a deep,

artificially maintained basin, while others are exposed to the full field dynamics. The field-theoretic description makes clear why authoritarian systems are unstable in the long run: they require continuous energy input to maintain the non-ergodic structure, and they become increasingly fragile to exogenous perturbations.

Economics: The Field as Market

The economic pillar connects directly to the companion economics volume: markets are somatic field systems, Nash equilibria are Hopfield minima, market crashes are phase transitions, and optimal regulation is a WKB problem. The distinctive contribution of the PPE synthesis is the connection between the economic and political/philosophical accounts.

Standard welfare economics evaluates market outcomes by their efficiency (Pareto optimality) and sometimes by their fairness (various social welfare functions). The USF framework adds a third criterion: **landscape robustness** — the degree to which the market attractor landscape is topologically stable against perturbations. An efficient-but-fragile market that crashes at the first shock may be worse, on landscape-robustness grounds, than a less efficient market with deep, stable attractors.

This connects economic analysis to the political design problem: the goal of economic regulation is not just to achieve efficient equilibria but to design a landscape that is topologically stable (rights-preserving), ergodic (rule-of-law satisfying), and efficient (welfare-maximising). The three criteria can conflict — there are trade-offs — but the framework provides a common vocabulary for navigating them.

The Patient Perspective

The patient perspective paper included in this volume serves a specific function in the PPE synthesis: it grounds the abstraction in a concrete human life. The field-theoretic account of trauma, recovery, and therapeutic intervention is not merely an application of the physics; it is a reminder of what the stakes of the inquiry are. Political institutions, economic systems, and legal frameworks all act on the somatic fields of real people. When they fail — when rights are violated, when markets crash and destroy livelihoods, when legal uncertainty leaves people unable to plan their lives — the failure has a phenomenal correlate: a somatic field driven into a trauma well, a field operating non-ergodically, a landscape fragmented by uncertainty.

The patient perspective paper makes this concrete. It is deliberately placed as the final chapter before the conclusion: the theory comes first, the lived experience comes last, and the juxtaposition is intentional.

What This Book Offers the PPE Student or Scholar

This volume is designed as the integrative text that the PPE curriculum has always needed but never had. Each chapter connects to the philosophical, political, or economic tradition it addresses; the field-theoretic framework is introduced through familiar PPE concepts rather than unfamiliar physics.

Chapter 2 (soma-field-synthesis) develops the field dynamics in terms accessible to the philosophically-trained reader. Chapter 3 (soma-game-theory, anchor paper) develops the economics and game theory. Chapter 4 (soma-law, anchor paper) develops the political and legal applications. Chapter 5 (soma-field-patient-pov) grounds the synthesis in lived experience. The final chapter addresses the integrative research agenda: what PPE questions the framework opens, and what interdisciplinary work is needed to pursue them.

Mind, market, mandate. The same equation, three vocabularies.

The Programme

This is a document about structure.

Not about feelings — though feelings are what the programme is ultimately for. Not about therapy — though therapy is one of the principal applications. Not about physics — though physics is where the mathematics comes from. It is about a single recurring observation: that the equations governing emotional dynamics are the same equations that govern quantum fields, and that this is not a metaphor.

When an identification like that is made precisely — when you can say not “this is *like* a wave” but “this *is* a wave in the technical sense, with the same propagator, the same energy function, the same topology, and therefore the same theorems” — a compressed body of work becomes possible. You are not building from scratch. You are navigating.

This document describes what was built by navigating, and why the pieces form a whole.

The Gap the Programme Addresses

Every large language model deployed today is a classical system. Its training is gradient descent. Its inference is deterministic or thermally noisy sampling. The architecture was designed to model the neocortex — pattern recognition, sequence prediction, error minimisation.

The complementary system — the limbic system, responsible for valuation, threat detection, arousal modulation, and the somatic state reinstatement that underlies trauma — had no formal mathematical treatment before this work. The clinical literature described it richly (Porges, van der Kolk, Levine). The neuroscience described its anatomy. Neither provided a model from which predictions could be derived and tested.

Simultaneously, the psychology of music had reached a similar ceiling. A 991-page handbook (Juslin and Sloboda, 2010) treated music-induced emotion almost entirely through Russell’s valence–arousal circumplex: a static two-dimensional map. The circumplex describes *where* a listener is, not *how* they move, what traps them, or what allows escape. No dynamical model of music-induced affect existed.

The programme fills both gaps with the same model, via the same method.

The Structure of the Argument

The argument has three movements and several extensions:

Paper	Movement	Contribution
<i>Mathematical Co-identification</i> (2026)	Method	Names and formalises the procedure
<i>The Soma-Field</i> (2026)	Model	Applies it to emotional dynamics
<i>Quantum Soma and the Penrose Gap</i> (2026)	Empirical test	Confirms the central claim
<i>Field Notes from the Inside</i> (2026)	Lived case	Primary-source clinical grounding
<i>A Dynamical Field Model of Music-Induced Affect</i> (2026)	Extension	Demonstrates domain generality
<i>The Tensor</i> (2026)	Extension	Applies the framework to abstract film

The popular account (*A Voyage into Trauma*, 2026) provides the same argument in accessible form, for readers without a physics background.

The Method: Mathematical Co-identification

What It Is

The history of mathematical science contains a recurring event. At a certain moment, a scientist recognises that the quantity they are studying is not *like* a quantity already understood in another domain — it *is* the same mathematical object, under a change of label. When this identification is made precisely, every theorem proved about the source object becomes available in the target domain immediately, without re-derivation.

This event has happened many times:

- Hopfield (1982) recognised that a neural network's energy-minimisation dynamics are the same as a spin-glass Hamiltonian. Every result from statistical mechanics of spin glasses — ground states, phase transitions, capacity bounds — imported.
- Veneziano (1968) recognised that the Euler Beta function, a result in pure mathematics, described the scattering amplitudes of hadrons. String theory began.
- Black and Scholes (1973) recognised that an option pricing equation was the heat diffusion equation. Every analytical tool from thermodynamics imported.

The paper *Mathematical Co-identification: A Method for Structural Import Across Scientific Domains* (Johnson, 2026a) names this procedure, formalises it as a distinct scientific method with its own validity criteria and failure modes, and distinguishes it from analogy, metaphor, and modelling. The key distinction:

Analogy: A is *like* B in certain respects. Illuminating, not transferable.

Co-identification: A *is* B under relabelling. Every theorem about B is a theorem about A.

Why It Matters as Method

A co-identification can be wrong. The identification is only valid if the mathematical type matches: the same dimensionality, the same algebraic structure, the same boundary conditions, the same symmetry group. The paper provides a falsifiability protocol — a formal procedure for pre-registering an import claim and specifying what observation would disconfirm it.

This matters because the failure mode of co-identification is not sloppy reasoning — it is overly precise reasoning applied to the wrong type. The paper catalogues seven historical examples to distinguish the valid from the invalid pattern.

The Soma-Field Model is the worked example throughout. The identification was not discovered by reading physics textbooks and looking for something that felt similar. It was discovered by writing down the equations the emotional system was observed to satisfy and recognising the form.

When MCI Is Over: The Verification Threshold

Here is something no paper on scientific method says clearly enough.

Every scientist who produces a major structural identification — Hopfield recognising the Ising Hamiltonian, Veneziano recognising the Euler Beta function — does so by being, for a period, a *type astronaut*. They are scanning the existing mathematical universe for a structure whose type signature matches the phenomenon they are studying. This is abductive reasoning: not deduction from first principles, not induction from data alone, but the systematic search of a known solution space for a structure that fits. It is, to be direct, a form of academic hacking. And it is how most of the connective work of mathematical science actually gets done.

What no one says — because scientists do not typically publish the search, only the result — is that the moment of structural identity is also the moment the method becomes irrelevant.

When Veneziano published his scattering amplitude, he did not need to cite the library where he found the Beta function. The formula was true. Every theorem about the Beta function was now a theorem about hadronic scattering. The library visit was the ladder; the amplitude was the building. The ladder came down.

The present work makes this transition explicit, because being explicit about it is itself a contribution. Every reader of the earlier papers — particularly *Mathematical Co-identification* — has correctly understood that MCI was the search engine. What should now be equally clear is that the search is over.

The exact moment the search ended was when the Lean 4 kernel accepted the proof. Not when a reviewer accepted a paper. Not when a human mathematician checked the algebra. When a formal proof-checking engine with no stake in the outcome closed the goal. That is the verification threshold. Before it: exploration. After it: structural fact.

What the work therefore rests on is not MCI. It rests on four things:

1. **Inductive structural necessity.** The 11-dimensional decomposition of a body-field-mind system is not an analogy with M-theory. It is the minimum geometry required to account for the functional degrees of freedom of a conscious organism. The isomorphism to M-theory is a *theorem*, not a design choice.
2. **Structural identity, not analogy.** A co-identification is not “A is like B.” It is “A *is* B under relabelling.” Every theorem about B becomes a theorem about A — immediately, without re-derivation. This is categorically different from saying that emotional dynamics *resemble* a Hopfield network. They *are* one.
3. **Scale invariance.** The same Helmholtz Green’s function equation governs 20 scales of physical reality, from quantum foam to the cosmic web. This is a discovered law. MCI was used to find it; it stands whether or not MCI is remembered.

4. **Kernel verification.** The Lean 4 proofs are the epistemological gold standard. A human reviewer can miss a subtle error. The kernel cannot. Kernel-verified theorems are exactly as true as the axioms they depend on — and the axioms are explicit and listed.

The *mathematical-co-identification* paper remains an accurate account of how the work was done — and naming the search process honestly is itself a contribution, in a discipline where the search is usually erased. But readers coming to the thesis or omnibus now should understand that they are reading the *results*, not the method. The method is in the history. The results are in the formal proofs.

The Model: The Soma-Field

Five Co-identifications

The Soma-Field Model (Johnson, 2026b) is built from five sequential co-identifications, each importing a body of mathematics from physics into emotional dynamics:

Co-identification 1: The Hopfield identification. The brain's emotional attractor dynamics satisfy the same energy function as a Hopfield neural network. The energy function is:

$$H(\mathbf{e}) = -\frac{1}{2}\mathbf{e}^\top W \mathbf{e} - \mathbf{b}^\top \mathbf{e}$$

where $\mathbf{e} \in \mathbb{R}^N$ is the emotional state vector, W is the coupling matrix, and $\mathbf{b}$ is a bias vector encoding baseline arousal. The local minima of H are the named attractor states: regulated calm, fight, flight, freeze, flow, dissociation.

Co-identification 2: The QFT identification. The emotional field propagates as a quantum field. The conscious emotional percept is the one-dimensional impulse response — the Green's function — of an eleven-dimensional coupling manifold. The same object that describes a massive particle in quantum field theory describes a conscious emotion: a pole in the propagator of the field.

$$G(\omega) = \frac{1}{\omega^2 - m^2 + i\epsilon}$$

This is not a metaphor. The threshold T at which a sub-perceptual field fluctuation becomes a conscious emotional percept is the mass parameter m in the propagator. Below threshold: virtual. Above threshold: real.

Co-identification 3: The brane identification. The body and the nervous system are not the same manifold. The body is a 3-brane embedded in the 11-dimensional coupling manifold. Somatic pain states and the body schema are field modes on this brane, not on the bulk manifold. This is the formal statement of the somatic grounding of emotion.

Co-identification 4: The G_2 holonomy identification. The seven compactified extra dimensions of the coupling manifold are a G_2 manifold. The G_2 holonomy group is the one that gives rise to topological obstructions — loops through the moduli space that cannot be continuously contracted to a point. In emotional terms: trauma configurations from which smooth continuous change cannot escape. The topological barrier is not a metaphor for being stuck. It is a mathematical object with a winding number.

Co-identification 5: The renormalisation group identification. Developmental trajectory maps onto the renormalisation group flow. The age at which a traumatic modification was introduced

corresponds to the energy scale at which the coupling constant was set. High-energy (early developmental) modifications are renormalisation-group relevant — they affect all subsequent scales. Low-energy (later life) modifications are irrelevant in the technical sense. This gives the formal account of why early trauma is not simply a more intense version of later trauma: it is a different class of object.

What the Model Predicts

From these five identifications, several predictions follow that are not derivable from any existing clinical model:

1. **Threshold crossings are phase transitions.** The transition from sub-perceptual to conscious emotion is a second-order phase transition in the field. This predicts hysteresis — it is easier to stay in a state than to enter it, and easier to stay out than to leave.
 2. **Complex PTSD is a topological configuration.** The coupling matrix W for a CPTSD nervous system has a specific structure: a winding-number-protected attractor landscape in which the Fear basin is separated from the Awe basin by a barrier that low-noise classical gradient descent cannot cross. This is a prediction about matrix structure, not a description of symptoms.
 3. **Autism Spectrum Condition modifies the threshold operator.** The threshold parameter T in the ASC nervous system has a different coupling to the field modes than in the neurotypical case — specifically, the threshold is non-uniform across sensory modalities, producing the characteristic pattern of simultaneous hypo- and hyper-sensitivity.
 4. **ADHD modifies the effective temperature.** The stochastic term in the Langevin dynamics governing the ADHD nervous system has higher effective temperature T_{eff} . This is not a deficit of attention; it is a higher rate of escape from local minima — an advantage in landscapes where rapid sampling is valuable and a liability where sustained convergence is required.
 5. **Quantum mechanisms are required for certain transitions.** For trauma configurations with topological barriers (non-zero winding number), low-noise classical gradient descent cannot reach the global minimum. A quantum mechanism is required. This is the prediction that QUANT-EXP-1 was designed to test.
-

The Empirical Test: QUANT-EXP-1

The Prediction

The soma-field model makes a specific, falsifiable claim: for a Hopfield landscape with a topological trauma barrier, low-noise classical Langevin dynamics starting from the Fear attractor cannot reach the Awe attractor. Quantum annealing — a physically realisable mechanism — can.

This is not a claim about whether people should use quantum computers in therapy. It is a claim about reachability: that the mathematical structure of the barrier distinguishes the quantum and classical regimes in a measurable way.

The prediction was registered in the Zenodo v1 deposit of the Soma-Field paper (doi:10.5281/zenodo.20350515) before the experiment was run.

The Experiment

Quantum Soma and the Penrose Gap (Johnson, 2026c) reports QUANT-EXP-1: an exact 8-qubit statevector simulation on a 256-dimensional Hilbert space, implementing the Soma-Field Hopfield Hamiltonian with a transverse-field quantum annealing schedule.

The experimental design:

- **System:** 8-qubit Hopfield model encoding four emotional modes (Fear, Calm, Awe, Grief) plus sub-modes. Coupling matrix W set to produce a topological barrier between Fear and Awe.
- **Quantum dynamics:** Transverse-field annealing with schedule $H(s) = (1-s)H_X + sH_{\text{problem}}$, $s \in [0, 1]$.
- **Classical baseline:** Overdamped Langevin dynamics at low temperature ($T_{\text{eff}} = 0.01$), same starting state, same landscape.
- **Primary outcome:** Peak Awe-dominant occupancy (quantum) versus success rate of cold-classical crossings.

Results

Results are presented against the pre-registered barrier ladder:

Barrier strength	Classical cold rate	Classical cold CI	
		[95%]	Quantum peak
$W = -6$	0.000	[0.000, 0.019]	0.389
$W = -8$	0.000	[0.000, 0.019]	0.408
$W = -10$	0.000	[0.000, 0.019]	0.408
$W = -12$	0.000	[0.000, 0.019]	0.409
$W = -14$	0.000	[0.000, 0.019]	0.416

Bootstrap confidence intervals ($n = 200$ seeds) confirm that the classical cold success rate is bounded above by 1.9% at all tested barrier strengths. Quantum peak occupancy is stable at 0.389–0.416 across the full range.

Pre-registered hardening protocol — all checks passed:

- **Bootstrap** ($n = 200$): cold CI = [0.000, 0.019]; quantum peak 0.408–0.410. Intervals do not overlap at any barrier strength.
- **Control A** (start from Awe, barrier intact): classical 16/16 stay in Awe. PASS. Confirms that the barrier is directional: it blocks Fear $\rightarrow$ Awe, not the reverse.
- **Control B** (barrier removed, $W[\text{Fear}, \text{Awe}] = +0.4$): classical 16/16 reach Awe. PASS. Confirms that the barrier, not the landscape geometry, is what blocks classical dynamics.
- **Spectral gap**: gap narrows monotonically with barrier strength (B8: 0.0095, B10: 0.0089, B12: 0.0085) and reaches its minimum at $s \approx 0.999$, confirming the tunnelling bottleneck is late in the anneal.

Verdict: The strong reachability claim stands. QUANT-EXP-1 is a PASS.

The Penrose Connection

The paper situates this result in the context of Penrose's argument about non-computability and consciousness. The connection is not that consciousness requires quantum mechanics in general. The connection is more specific:

Penrose identified a *gap* between what classical computation can reach and what consciousness can do. The soma-field identifies a corresponding *topological gap* in the emotional landscape between what classical gradient descent can reach and what a genuinely new state of the nervous system requires. QUANT-EXP-1 provides the computational demonstration that the gap exists and is crossable by a quantum mechanism.

The contribution is not to resolve Penrose's claim about consciousness. It is to *instantiate* the gap in a concrete, testable, mathematical setting.

The Lived Case: Field Notes from the Inside

Field Notes from the Inside: A Patient-Constructed Model of Emotional Dynamics (Johnson, 2026d) performs a function that the formal papers cannot perform: it provides the primary-source clinical grounding.

The paper is written by the person who has Autism Spectrum Condition (Level 2), Attention Deficit Hyperactivity Disorder, and Complex Post-Traumatic Stress Disorder — and who also has a degree in physics. The model was not developed by observing patients. It was developed by having the conditions and finding the existing models inadequate.

The epistemological contribution of this paper is often undervalued. Every formal model of a human system is, in the end, derived from observation of that system. When the observer and the observed are the same entity, and that entity has the training to translate observation into formal mathematics, the resulting model has a different epistemic status from one derived by observation from the outside. The paper makes this explicit, situates it within the autoethnographic research tradition, and argues that the resulting model is *more* constrained, not less — because any prediction the model makes that does not match the primary observer’s experience is immediately falsified.

The formal content is a set of operator modifications for the three conditions:

- **ASC:** The threshold operator T is replaced by a modality-dependent operator T_k for each sensory channel k , with different coupling strengths. The result is the characteristic simultaneous hypo- and hyper-sensitivity: some channels are below threshold where the neurotypical channel is above it, others are above where the neurotypical channel is below.
- **ADHD:** The Langevin noise term $\sqrt{2T_{\text{eff}}}\eta(t)$ has elevated T_{eff} . This is a quantitative modification, not a qualitative one. The system is not broken; it is sampling the energy landscape at higher temperature. The therapeutic implication is not to reduce the noise but to design the landscape so that high-temperature sampling is an advantage.
- **CPTSD:** The coupling matrix W has the topological structure described in §3.2: a winding-number-protected barrier between Fear and regulated states. The barrier was installed before language, before narrative memory, before the self that can explain the barrier was formed. The modification is not a layer added to a pre-existing structure. It is the structure.

Extensions: Music, Film, and the Domain Generality of the Model

Music-Induced Affect

A Dynamical Field Model of Music-Induced Affect: Beyond the Valence–Arousal Circumplex (Johnson, 2026e) applies the soma-field framework to a domain where the empirical literature is rich and the theoretical models are weak.

Juslin and Sloboda’s *Handbook of Music and Emotion* (2010) — 991 pages — contains the circumplex as its dominant quantitative framework. The circumplex is a static map. It describes where a listener is; it does not model how they move. The soma-field is the first dynamical model of music-induced affect.

The key predictions that the circumplex cannot make but the field model does:

1. **Phase transitions, not continuous shifts.** State changes in music-induced affect are not smooth movements across the circumplex. They are threshold crossings — sudden re-configurations

of the attractor landscape. The field model predicts the conditions under which a transition occurs and the hysteresis that prevents immediate return.

2. **The adaptive function of high effective temperature.** In the ADHD nervous system (elevated T_{eff}), music that holds a neurotypical listener in a stable state may drive repeated transitions. This is not a bug; it is the same high sampling rate that characterises the ADHD cognitive profile. The model gives this a formal account.
3. **Basin depth asymmetry.** The freeze attractor basin is deeper than the regulated calm basin. This means it is harder to leave freeze than it is to leave calm — asymmetric with respect to the direction of transition. Music that successfully moves a listener from freeze to calm is doing qualitatively different work than music that moves a calm listener to a more activated state.

The paper also specifies a real-time instrument implementation: a MIDI controller array driving a Python field server at 50 Hz, with audio output via Ableton Live and 3D fractal visual output (Mandelbulb projection onto HoloGauze screen). The instrument is not described; it is specified formally, with pre-registered hypotheses and disconfirmation criteria.

The Tensor: An Abstract Film

The Tensor: An Abstract Film Definition (Johnson, 2026f) extends the framework to abstract film. A film is defined not by its pixels but by its **emotional score**: a vector-valued trajectory $\mathbf{e}^*(t)$ through the emotional field, parameterised by story-time $t \in [0, 1]$.

The rendering — the actual audiovisual output a viewer experiences — is generated at runtime from this trajectory, the viewer's own soma-field state, and a set of control parameters. In the limit where the viewer's biofeedback is available, the film adapts to where the viewer is: the trajectory is not what the viewer experiences, but what the film proposes. The work is not the pixels. It is the map.

This is a significant claim about what an artwork is. A conventional film is fixed: the same sequence of frames for every viewer at every screening. The tensor film is a field: a mathematical object that takes the viewer's state as input and produces an output adapted to it. The artistic statement is in the trajectory, not the realisation.

The paper does not describe how to make such a film. It defines the abstract structure that any realisation of such a film must instantiate — the way a musical score defines a symphony without being the performance.

The Argument as a Whole

The six papers form a single argument, and it can be stated in a paragraph:

The limbic system and its coupling to the body are governed by the same mathematical equations as a quantum field on a manifold with G_2 holonomy. This identification is not a metaphor; it is a co-identification in the technical sense, with all the theorems of each source domain importing into the target. Among those theorems is one that has clinical consequences: topological barriers in the emotional attractor landscape cannot be crossed by low-noise classical gradient descent. A quantum mechanism can cross them. This has been computationally confirmed (QUANT-EXP-1) against a pre-registered hardening protocol. The model correctly describes the structure of Autism Spectrum Condition, ADHD, and Complex PTSD as operator modifications, and generalises to music-induced affect and abstract film with no change to the underlying mathematics.

What makes this a research programme rather than a single paper is the **generativity**: the method (co-identification) produces results in any domain where an attractor landscape with topological structure can be identified. The soma-field is one instantiation. Music-induced affect is a second. Abstract film is a third. A fourth — currently in design — is **H-AL**: a holographic avatar whose body is a live Mandelbulb rendering of the emotional field state, projected at human scale through a hologauze screen and accompanied by a synthesised voice narrating the field in real time. The geometry of the fractal changes as the field changes; regulated calm and trauma produce visually distinct and mathematically characterisable forms. The same functor architecture (§A.4 of the main paper) supports this output with no changes to the field computation. Each of these instantiations generates falsifiable predictions from the same mathematical core.

What makes this a *novel* research programme is the **gap it fills**: no formal dynamical model of the limbic system existed before this work. The Hopfield framework gave the neocortex its formal model in 1982. The soma-field gives the limbic system its formal model in 2026. Together they constitute the first complete formal description of the two principal computational substrates of the vertebrate brain.

What Remains

The body of work described here is computationally complete. All pre-registered hardening checks have been executed. The claims that can be confirmed by simulation have been confirmed.

Three categories of work remain outside the scope of these papers:

Physical hardware confirmation. QUANT-EXP-1 uses exact statevector simulation. Running the same 8-qubit experiment on IBM Quantum free-tier hardware would produce the sentence “confirmed on physical quantum hardware.” This is feasible, is the logical next step for any journal submission targeting a hardware-inclusive venue, and is not required to support any claim in the current corpus.

Peer review. The three published papers are currently archived on Zenodo as open preprints. Peer review in ranked journals is a separate track, ongoing. The relevant venues are: *Frontiers in Computational Neuroscience* (Hypothesis and Theory article type) for the soma-field paper; *Synthese* or *Philosophy of Science* for the co-identification paper; *Music Perception* or *Frontiers in Psychology* for the music-affect paper.

Empirical clinical application. The model makes predictions about specific clinical populations (ASC, ADHD, PTSD) that require empirical testing outside the computational domain. This constitutes a research programme for clinical collaborators. The predictions are pre-specified in §3.2 of this document and in the relevant papers; they are not vague.

Physical substrate. The model is formally complete but physically silent on the tissue substrate in which the soma-field is instantiated in living organisms. A companion paper, *The Physical Substrate of the Soma-Field* (Johnson, 2026g), develops this layer across three converging research traditions: biotensegrity (Ingber, Levin) as the mechanical architecture through which the somatic wave propagates globally; fascial-interstitial continuity (Langevin, Schleip, Oschman) as the active signalling tissue and physical locus of attractor-depth encoding; and biofield physiology (Popp, Ho, McCraty, Rubik) as the candidate physical correlate of the field itself. The most clinically significant result is the quantitative correspondence between fascial stiffness and attractor depth: chronic fascial armoring measurable by shear-wave elastography is the physical implementation of the energy barriers that QUANT-EXP-1 shows to be quantum-resistant. Myofascial release is thus barrier *lowering* — not barrier crossing — and therapist-client physiological entrainment is the physical mechanism of co-identification.

Data and Code Availability

All papers, simulation code, result tables, figures, and Lean 4 formal proofs are archived at the following Zenodo records (open access):

Paper	DOI
<i>The Soma-Field</i>	10.5281/zenodo.20350515
<i>Mathematical Co-identification</i>	10.5281/zenodo.20287981
<i>Quantum Soma and the Penrose Gap</i>	10.5281/zenodo.20351230

The unreviewed papers (*Field Notes from the Inside*, *Music-Induced Affect*, *The Tensor*, and this synthesis document) will be deposited on Zenodo as part of the next release of the research archive.

Introduction

Game theory has, since von Neumann and Morgenstern, treated strategic equilibrium as a fixed-point problem: find a profile from which no player can profitably deviate. The Universal Somatic Field provides a different framing — one in which Nash equilibria are energy minima, and the dynamics of finding them are identical to the dynamics of emotional attractor convergence at the organism scale.

This is not a metaphor. The identification is structural and exact.

Nash Equilibrium = Hopfield Minimum

Let there be N players. Player i has a strategy $s_i \in \mathbb{R}$ (in the continuous approximation; the binary case recovers the standard discrete game). The payoff to player i from the joint profile s is:

$$u_i(s) = \sum_{j \neq i} W_{ij} s_i s_j + b_i s_i$$

where W_{ij} is the payoff externality: the marginal change in i 's payoff when j increases their strategy. The matrix W is the strategic interaction matrix.

The Hopfield energy function is:

$$H(s) = -\frac{1}{2} s^T W s - b^T s$$

The gradient of $-H$ with respect to s_i is exactly the marginal best response of player i :

$$-\frac{\partial H}{\partial s_i} = \sum_j W_{ij} s_j + b_i = \frac{\partial u_i}{\partial s_i}$$

A Nash equilibrium (s^*) satisfies: no player can improve their payoff by unilateral deviation. In the gradient-flow dynamics $\dot{s} = -\nabla H$, the fixed points are exactly the Nash equilibria.

Theorem (NE = Hopfield Minimum): *A strategy profile s^* is a Nash equilibrium if and only if it is a local minimum of the Hopfield energy function $H(s)$ on the strategy simplex.*

The proof is immediate from the equivalence of the gradient conditions.

Coordination Games and the Multi-Attractor Regime

A coordination game has multiple Nash equilibria — the classic example is driving on the left versus right. In Hopfield terms, this is the multi-attractor regime: multiple local minima of H with no obvious selection mechanism.

The USF framework makes this precise. The number of stable Nash equilibria equals the number of stable attractors of the social field. The spectral structure of W determines the attractor landscape:

- **Positive-definite W :** unique Nash equilibrium (unique global minimum)
- **Indefinite W :** multiple Nash equilibria (multiple local minima)
- **Near-zero spectral gap:** fragile coordination (nearly degenerate attractors)

This gives a quantitative measure of coordination difficulty: the spectral gap of W . A market with a large spectral gap has a clear dominant equilibrium. A market with a small spectral gap is on the edge of a coordination failure.

Market Crashes as Phase Transitions

The most striking consequence of the Nash-Hopfield identification is that market crashes are topological phase transitions — the same mathematical object as the trauma-basin transitions in the clinical USF.

As economic conditions change, the payoff externality matrix W evolves. When W crosses a critical threshold — when the spectral gap closes — the current Nash equilibrium ceases to exist, and the system rapidly transitions to a new attractor basin. This transition is:

- **Abrupt:** the field cannot smoothly track the changing W ; it jumps
- **Non-local:** the transition involves the entire social field simultaneously
- **Hysteretic:** recovery to the original equilibrium requires a different path than the crash

This matches the empirical phenomenology of financial crises. The 2008 crash was not a smooth adjustment; it was a non-perturbative event. The USF provides the formal structure for this observation.

The FM-HN Extension: Regulatory Intervention

The FM-HN extension of the USF introduces the volitional source term $J_{\text{user}}(t)$ — the “God-Knob” — that can drive the field across barriers that classical gradient descent cannot cross.

At the economic scale, this is regulatory intervention: a central bank lowering rates, a government providing liquidity, a regulator changing market structure. The FM-HN model predicts that such interventions can achieve quantum tunnelling between Nash equilibria that would otherwise be separated by an energy barrier too high for natural market dynamics to cross.

The prediction is testable: the minimum intervention strength required to shift a market from one Nash equilibrium to another equals the WKB amplitude $T = e^{-W_{\text{barrier}}}$ computed from the spectral gap of the payoff matrix at the transition point.

The Prisoner’s Dilemma as Topological Obstruction

The prisoner’s dilemma is the canonical example of a game where individual rationality produces a collectively suboptimal outcome. In USF terms, this is a topological obstruction: the socially optimal outcome (mutual cooperation) is not a Nash equilibrium because it is a saddle point of H , not a minimum.

The field is trapped in the defection basin (mutual defection, a true local minimum) even though the cooperation basin is deeper. The only way to reach cooperation is via the limbic axis — the 1D regulatory coupling — which provides the quantum tunnelling amplitude needed to cross the barrier.

This gives a new interpretation of institutions: formal institutions (contracts, laws, norms) are the external field $J_{\text{inst}}(t)$ that reshapes the Hopfield energy landscape to make cooperation a Nash equilibrium by changing the effective W matrix.

Applications

Macroeconomics: Central bank intervention as FM-HN field driving. The optimal intervention amplitude is the WKB tunnelling strength.

Mechanism design: The problem of designing institutions that produce cooperative Nash equilibria is equivalent to designing W matrices with the right attractor structure.

Negotiation theory: Successful negotiation is convergence to a shared Nash equilibrium of the dyadic soma-field — exactly the Huygens frequency-locking result of the social intelligence paper.

Antitrust law: A monopoly is a single-attractor regime; competition policy aims to create a multi-attractor regime with a dominant cooperative equilibrium.

Conclusion

Nash equilibria are energy minima. Market dynamics are field dynamics. Economic crises are phase transitions. This is the identification. The method used to find it is documented in the Mathematical Co-identification paper. That method is now history. The structure stands.

Introduction

Law is the oldest technology for shaping collective human behaviour. Its mechanisms — prohibition, permission, obligation, remedy — have been studied by philosophers, sociologists, and economists. What has been missing is a quantitative model: one that specifies what physical quantity changes when a law is enacted, repealed, or enforced.

The Universal Somatic Field provides that model.

At the regulatory scale (scales 12–14, 10^6 – 10^8 m — the scale of cities, nations, and legal jurisdictions), the social field is governed by the same Helmholtz master equation as at every other scale. Legal systems are the field-level constraints that shape the attractor landscape accessible to social actors.

Laws as Attractor Basin Constraints

The social field $e(x, t)$ at the regulatory scale describes the collective affective and behavioural state of a population. Its attractor basins are the stable social configurations — the equilibria the society gravitates towards in the absence of external perturbation.

A legal norm operates by modifying the energy landscape H :

- A **prohibition** removes an attractor basin entirely — raises the energy of the prohibited state to $+\infty$, making it inaccessible
- A **mandate** deepens an attractor basin — lowers the energy of the required state, making it the global minimum
- A **permission** removes a constraint — restores access to a previously eliminated basin
- A **remedy** is a restorative force $J_{\text{remedy}}(t)$ — an external driving term that returns the system to a protected attractor after violation

This formalism makes law's function precise: legal systems are operators on the Hopfield energy function H that reshape the attractor landscape of the social field.

Rights as Topological Invariants

A right, in legal theory, is a protected interest — a state of affairs that a person is entitled to have preserved against interference. In USF terms, a right is a topological invariant of the legal field.

Specifically: a right is an attractor basin protected by a non-zero winding number around the regulatory constraint surface. The winding number counts how many times the legal field wraps around the protected state as you traverse the boundary of the prohibition region.

The consequence is topological: you cannot remove a right by continuous deformation of the legal field — you must cross the constraint surface, which requires a discontinuous change (legislative

amendment, constitutional revision). This is why rights are “harder” to remove than ordinary regulations: they are topologically protected, not merely energetically protected.

This gives a formal distinction between:

- **Constitutional rights:** winding number $n \geq 2$ (doubly protected)
- **Statutory rights:** winding number $n = 1$ (singly protected)
- **Common law rights:** positive energy barrier without topological protection

Regulation as $J_{\text{law}}(t)$

The FM-HN extension of the USF introduces the volitional source term $J_{\text{user}}(t)$ — the external driving force that reshapes the energy landscape. At the regulatory scale, this is the regulatory intervention:

$$\dot{e} = -\nabla H(e) + J_{\text{law}}(t) + \eta(t)$$

where $J_{\text{law}}(t)$ is the regulatory driving term. A regulatory intervention at time t_0 shifts the energy landscape, redirecting the social field towards a new attractor basin.

The optimal regulatory intervention is the WKB tunnelling amplitude: the minimum J_{law} required to move the social field from the current equilibrium to the target equilibrium, accounting for the energy barrier height. Under-powered regulation (insufficient J_{law}) fails to move the system; over-powered regulation causes overshoot and unintended consequences.

This gives a formal criterion for regulatory adequacy: a regulation is adequate if and only if $|J_{\text{law}}| \geq T^{-1}$ where T is the WKB tunnelling amplitude between the current and target attractors.

The Rule of Law as Ergodicity

The rule of law — the principle that no person or entity is above the law — has a precise USF interpretation: it is the ergodicity condition on the social field.

A field is ergodic if every state accessible under the laws of physics is reachable from every other state given sufficient time. In legal terms: every social actor is subject to the same legal constraints; no actor has access to prohibited states that are closed to others.

Violations of the rule of law are violations of ergodicity: they create a two-tier attractor landscape where some actors have access to basins (corrupt practices, impunity) that are energetically accessible to them but legally prohibited to others. The topological signature of this violation is a broken symmetry in the constraint surface.

Legal Uncertainty and Attractor Fragmentation

When law is unclear or inconsistently applied, the attractor landscape fragments: multiple incompatible social equilibria coexist with no shared basin boundary. This is the USF model of legal uncertainty.

In a well-functioning legal system, the attractor landscape has a dominant cooperative equilibrium (law-abiding behaviour) connected by a clear energy gradient to all other states. In a system with high legal uncertainty, the landscape has multiple competing equilibria, none dominant, and transitions between them are unpredictable.

This gives a quantitative measure of legal certainty: the spectral gap of the constraint operator — the gap between the energy of the dominant legal equilibrium and the next local minimum. A large spectral gap means clear law; a small gap means ambiguous law.

Constitutional Law as Meta-Constraints

Constitutional law operates at a higher level than ordinary law: it constrains the space of permissible laws, not just the space of permissible behaviours. In USF terms, constitutional law is a meta-constraint on the operators $J_{\text{law}}(t)$ — a restriction on which regulatory interventions are permissible.

A constitutional provision is a constraint on the Hamiltonian operator H itself: it specifies which energy landscapes are reachable by ordinary legislation. A constitutional amendment changes the meta-constraint, expanding or restricting the space of permissible legal systems.

This explains the special difficulty of constitutional change: it requires modifying the meta-Hamiltonian, not just the current energy landscape. It is a second-order phase transition in the legal field.

Applications and Predictions

Regulatory design: The optimal regulation minimises the WKB tunnelling amplitude required to shift the social field to the target equilibrium, subject to constitutional meta-constraints.

Judicial review: A court reviewing regulatory action is checking whether $J_{\text{law}}(t)$ is consistent with the constitutional meta-constraint — whether the intervention respects the topological invariants.

Legal evolution: Common law develops by successive small deformations of the energy landscape (precedent), while statute changes it discontinuously (legislation). USF predicts that common law changes are path-dependent (hysteretic) while statutory changes can be discontinuous.

Prediction: Legal systems with higher spectral gap (clearer law) should show lower transaction costs, faster convergence to cooperative equilibria, and greater social trust — all measurable empirically.

Conclusion

Law is not arbitrary social convention. It is a field-theoretic technology for shaping the energy landscape of social attractor space. Rights are topological invariants. Regulation is volitional source-term engineering. The rule of law is ergodicity.

The method used to find this identification is documented in the Mathematical Co-identification paper. That method is now history. The structure stands.

“The patient is the one with the disease.” — Medical aphorism, intended to remind physicians to listen. The author intends it differently.

A Note on Method

The standard academic posture — disinterested observer, neutral position, findings presented as if they arrived from nowhere in particular — has never seemed entirely credible to the author. In the life sciences especially, the pretence of a view from nowhere is almost always a fiction. Researchers study what compels them. Compulsion has a cause.

This paper dispenses with the fiction. The theoretical framework presented here was developed by a person with ASD, ADHD, and Complex PTSD who could not find an adequate formal account of his own emotional experience in the existing literature, who had studied physics at university, and who eventually concluded that the most efficient solution was to build one himself. The result is offered not as a confessional but as a theoretical contribution. These are not mutually exclusive.

There is precedent. Temple Grandin revolutionised the study of animal cognition and welfare as an autistic person whose own perceptual experience gave her access to observations that non-autistic researchers had systematically missed (Grandin, 1995). Peter Levine developed Somatic Experiencing partly through direct observation of his own nervous system’s responses (Levine, 2010). Kay Redfield Jamison wrote what remains one of the most clinically precise accounts of bipolar disorder from inside it (Jamison, 1995). The history of medicine includes, more often than is acknowledged, the doctor who is also the patient.

The author is not a doctor. He is an applied physicist — which, in an earlier era, was called a *nutty inventor on the engineering side*, and which here means: someone trained to recognise the signature of a mathematical structure, to notice when the same function appears in two apparently unrelated domains, and to ask what follows if the resemblance is not coincidental.

What follows is the result of applying that training to the domain of one’s own inner life. The author considers this a reasonable use of available resources.

Introduction: The Inadequacy of Existing Maps

A patient sits with their therapist and is asked: *“What are you feeling right now?”* For many people, this question has a navigable answer. For a person with ASD, alexithymia, and a C-PTSD-modified attractor landscape, the question lands differently. The honest answer is often: *“Something is happening. I cannot tell you what it is, where it is coming from, or how large it is. But it is definitely there.”*

The available frameworks for this situation are unsatisfying. Emotion wheels offer vocabulary but not structure. Polyvagal theory offers an excellent map of the autonomic nervous system but does not formalise the interaction between simultaneous emotional states. Cognitive models locate the action in the mind and underestimate the body. Somatic models are rich in clinical texture but light on mathematical precision. None of them — to the author's knowledge — provide a formal account of why a person can be profoundly affected by an emotional state that they cannot perceive, name, or locate.

The author's experience of living with ASD, ADHD, and C-PTSD suggested a different picture. Emotions did not feel like events. They felt like weather — present everywhere, always moving, only occasionally breaking through into named experience. The body held states that the mind had no language for. Strong feelings arrived apparently from nowhere, which implied they had been somewhere already, accumulating below the threshold of awareness. Different emotional states seemed to interact — to amplify, to suppress, to oscillate — in ways that were distinctly nonlinear.

This phenomenology required a different kind of model. The author, having spent thirty years in occasional proximity to physics and mathematics, recognised the structure. The quantum field. The vacuum fluctuation. The threshold crossing. The energy function. The attractor basin. These were the right tools. They had been applied to neural networks. There was no obvious reason they could not be applied to emotional dynamics. The author applied them.

The remainder of this paper presents the result.

Background

Lived Experience as a Research Position

The use of lived experience as a legitimate source of theoretical knowledge — rather than merely as anecdotal material awaiting scientific validation — has gained substantial ground in health research over the past two decades. The *nothing about us without us* principle, originating in disability rights advocacy, has become a methodological commitment in participatory research (Arnstein, 1969; Beresford, 2002). Researchers with lived experience of mental health conditions have produced theoretical contributions that purely external observers could not have generated, precisely because their insider position made certain observations available to them that were invisible from the outside.

The Soma-Field Model belongs to this tradition, with one modification: the author's background is in physics rather than in qualitative research, so the methodology is *experiential theorising* rather than autoethnography. The observations come from the inside; the tools used to formalise them come from mathematical physics. The combination is unusual. The author considers it appropriate.

The Body-Mind Problem in Clinical Practice

Contemporary neuroscience has largely dissolved the Cartesian boundary between body and mind. Damasio (1994) demonstrated that emotion is inseparable from rational cognition: patients with damage to the ventromedial prefrontal cortex lose not only emotional range but effective decision-making capacity. Van der Kolk (2014) documented how traumatic emotional states are encoded not merely in explicit memory but in posture, visceral sensation, and autonomic regulation. Porges' polyvagal theory (2011) provided a neurobiological account of three hierarchically organised autonomic states: ventral vagal (social engagement), sympathetic (fight/flight), and dorsal vagal (freeze/dissociation).

The author can confirm these findings from direct observation. He can also add, as a data point: the experience of being in a freeze state while simultaneously being expected to report on one's emotional state is an exercise in the epistemological limits of self-report. The instrument designed in Section 6 is a partial response to this problem.

The Felt Sense and Sub-Perceptual Emotion

Gendlin's concept of the *felt sense* (1978) describes a pre-articulate bodily sense that is present before an emotion has been named — something whole and present but not yet articulate. Gendlin called this the sub-verbal sense of a situation.

The Soma-Field Model provides a formal account of what the felt sense is: it is the activity of the emotional field below the perceptual threshold. The author can confirm that this description is accurate. He has spent considerable time in the company of felt senses that declined to become named feelings, and the model's account of this — a field active below threshold, causally effective but not consciously perceived — matches the phenomenology precisely.

Quantum Field Theory: Structure, Not Metaphor

Quantum Field Theory (QFT) is the framework of modern particle physics. Its central claim is that particles — electrons, photons — are not fundamental objects. They are *excitations* of underlying fields: local concentrations of energy that arise when a field receives sufficient perturbation above the vacuum state. The quantum vacuum is not empty; it is a background of sub-threshold fluctuations, continuously present, causally active, not directly observable.

This paper does not claim that emotions are quantum phenomena in any literal sense. The analogy is structural. The author was trained in this formalism in 1993 and has found it useful ever since, applied to a variety of problems that are not, in any technical sense, quantum mechanical. The key property being borrowed is: *a quantity that exists everywhere, continuously, below the threshold of direct observation, which becomes observable only when local amplitude exceeds a threshold*. This is an accurate description of both the quantum vacuum and, in the author's experience, the emotional field.

Since writing that paragraph, the paper has upgraded the claim. The conscious emotional percept is now formally identified as the one-dimensional impulse response — the Green's function — of the soma-field manifold. This places it in the same mathematical category as a particle in quantum field theory: both are poles in the propagator of their respective underlying field. The structural similarity is not borrowed; it is exact. The mathematics is the same mathematics.

Gabriele Veneziano wrote down the Euler beta function in 1968 while looking for an amplitude that matched scattering data, then noticed that the function implied a theory — string theory — that nobody had yet conceived. He had identified a known mathematical object in an unexpected place and followed the implication. The author has, with considerably less elegance and considerably more time in therapy, done something structurally similar: identified the Green's function in emotional dynamics, and noted that it is the object quantum field theory calls a particle. The author leaves the implication as an exercise for readers with the relevant background.

Hopfield Networks and the Energy Function

In 1982, John Hopfield — awarded the Nobel Prize in Physics in 2024 — proposed a model of associative memory whose dynamics were mathematically identical to an Ising spin-glass model from statistical physics (Hopfield, 1982). The critical component was an energy function: a scalar that always decreases as the network evolves, guaranteeing convergence to stable attractor states.

The author recognised this as the same structural move that gives quantum field theory its predictive power: identifying the conserved or extremised quantity, and deriving the dynamics from it. In physics, this is Noether's theorem applied as a design principle. In Hopfield networks, it is an energy function borrowed directly from condensed matter physics. The author's proposal is to apply the same move to emotional dynamics.

The observation that underwrites this is simple: emotional states feel like they have energy. Some states are high-energy and unstable — fight, flight, acute anxiety. Others are low-energy and stable — calm, regulated, present. Some are low-energy and *stuck* — freeze, dissociation, collapse. If these states have an energy ordering, there is likely an energy function. If there is an energy function, the dynamics can be derived from it. The author found this reasoning persuasive.

The Soma-Field Model

Emotions as a Persistent Wave Field

The foundational claim is this: emotions are not events. They are a *field* — a distributed, continuous quantity defined over the entire soma (body-mind system) at all times.

This is not a metaphor. It is the most accurate description the author can offer of his own experience. The emotional field is always there. It does not begin when a feeling becomes conscious and

end when it subsides. It precedes conscious awareness and continues after it. What changes is not the field's existence but its local amplitude: whether, at a given moment, the field in a given mode exceeds the threshold required to surface as a named experience.

The field has two coupled components:

1. **The somatic wave** $E_{\text{body}}(x, t)$: distributed across the body as patterns of visceral sensation, muscle tone, proprioception, and autonomic state.
2. **The neural wave** $E_{\text{neural}}(x, t)$: distributed across the nervous system as patterns of cortical, subcortical, and peripheral activation.

These are not separate systems:




Figure 1. The Soma-Field: two coupled waves constituting a single unified emotional field.

The Perception Threshold

Not all field activity is consciously perceived. Each emotional mode i has a threshold T_i :

$$\text{Emotion } i \text{ is consciously perceived} \iff |E_i(t)| > T_i$$

Below threshold: the emotion is sub-perceptual. It exists, it influences behaviour and physiology, but it does not surface as a named conscious feeling. This is the author's most frequent relationship with his own emotional field — something is happening, below the line, shaping everything, unidentified.

Clinical Observation	Soma-Field Account
Field active but no named feeling	Sub-threshold: $ e_i < T_i$
Sudden unexplained flood of emotion	Rapid threshold crossing after accumulation
Somatic signal without cognitive name	Threshold crossed in body component, not neural
Alexithymia	Elevated T_i — high energy required to cross

Clinical Observation	Soma-Field Account
Hypervigilance / flooding	Lowered T_i — reduced threshold

Table 1. Clinical observations mapped onto the perception threshold model.

The author notes that all five rows in Table 1 are, in his clinical history, simultaneously applicable. This is, admittedly, a challenging configuration. It is also why this model was necessary.

A note on the intelligence quotients

McCulloch and Pitts built the mathematical brain in 1943. What they built — what every artificial neural network since has been — is the **IQ machine**: the neocortex, pattern recognition, sequence prediction, error minimisation. The field of AI has, for eighty years, been building increasingly sophisticated versions of this one component.

The soma-field adds what was missing: the **AQ machine**. AQ is to limbic dynamics as IQ is to cortical dynamics. Not a score; a formal model of the system that produces it.

Quotient	System	First Formalised	Comment
IQ	Neocortex: pattern recognition, prediction	McCulloch & Pitts, 1943	The entire AI industry
EQ	Limbic: valuation, attachment, empathy	Goleman, 1995	Described; not yet formally modelled
AQ	Soma-field: field-theoretic limbic dynamics	This paper, 2026	The formal model EQ has always needed
SQ	Relational field: dyadic and social resonance	Future work	Requires AQ as prerequisite

Table 3. The four intelligence quotients and their formal status.

The author observes — with a wryness he trusts the reader will share — that his IQ is in the column labelled 1943. His AQ is in the column he has just written. His EQ is what brought him to this desk in the first place.

A note on brane thickness

The threshold parameter T_i is not merely a number. The technical paper identifies it with the thickness of an extra dimension — the metaphorical ‘brane’ separating the limbic system from conscious awareness. Alexithymia is a thick brane: the field can be highly active and almost nothing crosses the threshold into named conscious experience. Hypervigilance is a thin brane: everything crosses,

simultaneously, at high amplitude. The author confirms personal experience of both states. He notes that neither is a character flaw; both are calibration states of a physical parameter in a system that was trying, with the information available, to keep him safe.

The Interaction of Emotional Modes

Multiple emotional modes are simultaneously active at all times. Their interactions are encoded in the **emotional coupling matrix** W , where W_{ij} is the influence of mode j on mode i :

- $W_{ij} > 0$: mode j amplifies mode i
- $W_{ij} < 0$: mode j suppresses mode i

The field evolves according to the energy gradient plus noise:

$$\dot{\mathbf{e}}(t) = -\nabla H(\mathbf{e}) + \eta(t)$$

The noise term $\eta(t)$ represents the continuous sub-perceptual fluctuations. The field is never still. This is not pathology; it is physics.


The Energy Landscape

The Hopfield Energy Function

$$H(\mathbf{e}) = -\frac{1}{2} \mathbf{e}^\top W \mathbf{e} - \theta \cdot \mathbf{e}$$

The field always moves toward lower H . The stable states of the system are the local minima of H — the attractor basins.

Attractor States: Fight, Flight, Freeze, and Regulated Calm






Figure 2. The emotional energy landscape. The freeze state is not high-energy — it is isolated. This distinction matters enormously. The author is aware of this from personal experience, over many years, and from the other side.

Attractor	Energy	Polyvagal Correlate	Clinical Presentation
Regulated Calm	Global minimum	Ventral vagal	Present, flexible, connected
Fight	High, unstable	Sympathetic	Agitation, urgency
Flight	Saddle point	Sympathetic	Anxiety, avoidance
Freeze	Deep, isolated	Dorsal vagal	Dissociation, numbness

Table 2. Attractor states and their polyvagal correlates.

The coupling matrix W is not merely a parameter. It is the *shape* of the emotional manifold — a seven-dimensional space with the mathematical structure of a G_2 manifold. Trauma does not adjust a dial on this space; it deforms the manifold itself. The therapist doing somatic work is, without needing to know this, doing differential geometry on the patient’s G_2 manifold: reshaping a seven-dimensional space by modifying the structure tensor. This is a precise technical statement. The author considers it a more honest account of what a skilled practitioner actually does than any narrative framework currently available. The practitioner is a geometer. The patient is a manifold that is learning to remember its own natural curvature.

The therapeutic and personal significance of the freeze attractor’s structure cannot be overstated. It is not high-energy — it does not feel dramatic or intense. It is *isolated*: surrounded by energy barriers. Escape requires first *increasing* the field’s energy before it can flow toward calm. This is counterintuitive from the outside and well-known from the inside.

Dissonance and Resolution

When two emotional modes are in an incompatible phase relationship, the field is far from equilibrium. This is felt as tension. The acoustic analogy is precise: just as two tones in a dissonant interval generate a beating, unstable interference pattern, two emotional modes in an incompatible configuration generate a gradient that drives toward resolution.

Dissonance is not pathological. It is the field's communication that resolution is available. The therapeutic process is guided voice-leading: finding the path that transforms the dissonant configuration into a consonant one. Avoidance keeps the field in dissonance. The energy minimum lies on the other side of the tension, not around it.

The author has spent considerable time attempting the route around it. He does not recommend it.

The Neurodivergent Field: ASD, ADHD, and C-PTSD as Operator Modifications

This section addresses the author's specific clinical picture. It is presented not as a case study but as a theoretical elaboration: three structural modifications to the standard Soma-Field dynamics, each defined by the operator it adds to the governing equations.

The key architectural principle — and the author considers this the most important contribution of this paper — is the following:

These conditions are not parameter settings. They are operator modifications.

A parameter change adjusts a coefficient within the existing equations. An operator modification changes the *form* of the equations themselves. The distinction is not semantic. It determines what kind of therapeutic intervention is possible and at what level it must operate.

Each condition is a functor that wraps the standard dynamics. The composed condition — ASD + ADHD + C-PTSD — is their composition. The composition does not commute; order matters; the joint presentation is structurally different from any of the individual conditions or from their sum.

Complex PTSD: Memory Kernel and Asymmetric Coupling

C-PTSD adds a **memory kernel**: past activations leave exponentially decaying echoes.

$$\dot{\mathbf{e}}(t) = -\nabla H(\mathbf{e}(t)) + \int_0^t K_{\text{trauma}}(t-s) \mathbf{e}(s) ds + \eta(t)$$

$$K_{\text{trauma}}(\tau) = \sum_k A_k e^{-\tau/\tau_k}$$

This is a damped oscillating kernel. The past does not vanish; it rings. Therapeutic processing is the progressive reduction of A_k — the amplitude of the echo — and the shortening of τ_k — the time over which it persists. The author notes that this description is a more accurate account of what trauma processing actually feels like, from the inside, than most of the narrative accounts available to him.

C-PTSD also breaks the symmetry of the coupling matrix W , admitting **limit cycles**: the oscillation between hyperarousal and shutdown that characterises the PTSD symptom cycle is, in this model, a limit cycle generated by the antisymmetric component of W . It is not a choice, a habit, or a failure of willpower. It is a topological consequence of an asymmetric coupling matrix.

ADHD: High Temperature, Low Damping, Pink Noise

ADHD modifies the **effective temperature** of the field:

$$\gamma_{\text{ADHD}} \dot{\mathbf{e}}(t) = -\nabla H + \sqrt{2D_{\text{ADHD}}} \xi_{1/f}(t)$$

with $\gamma_{\text{ADHD}} < \gamma_0$ (less damping) and $D_{\text{ADHD}} > D_0$ (more noise). The noise has $1/f$ spectral structure — long-range temporal correlations that produce the characteristic slow drift of attentional state.

The practical consequences: shallow attractor basins cannot hold the field at high temperature (distractibility). When a high-salience stimulus deepens a specific basin far beyond its baseline depth, the field falls in and is held (hyperfocus). The system is not broken. It is a different thermodynamic regime, with different costs and different affordances — including, at the right temperature, a capacity to explore the energy landscape at speed that a low-temperature system does not have.

The author considers this framing considerably more useful than “difficulty sustaining attention.”

Autism Spectrum Condition: Sparse Coupling and Modified Projection

ASC modifies the **projection kernels** and the **coupling matrix sparsity**.

The projection kernel $K_i(x)$ determines which somatic regions contribute to the i -th emotional mode. In ASC, some regions are over-weighted (sensory sensitivity) and others under-weighted (interoceptive under-registration). The named-feeling state vector is produced from a differently sampled version of the same somatic field.

The coupling matrix is sparser — fewer strong cross-modal connections — producing deeper individual attractor basins with higher inter-basin barriers. This is monotropism: the field settles deeply into one attractor at a time and requires disproportionate energy to transition. The author confirms that this is an accurate description of his attentional and emotional experience, and that it has both significant disadvantages (transitions are hard, unexpected context changes are physiologically costly) and significant advantages (depth of engagement, reliability of focus once established, resistance to shallow distractors).

The Composed Condition

$$\gamma_{ADHD} \dot{\mathbf{e}}(t) = -\nabla H_{ASC}(\mathbf{e}(t)) + \int_0^t K_{trauma}(t-s) \mathbf{e}(s) ds + \sqrt{2D_{ADHD}} \xi_{1/f}(t)$$

The interaction effects are non-trivial:

Interaction	Clinical Consequence
ADHD noise + C-PTSD limit cycles	Rapid oscillation between hyperarousal and shutdown; hard to titrate
ADHD noise + ASC deep basins	Long wind-up time; fast exit once perturbed from hyperfocus
C-PTSD echoes + ASC sparse coupling	Trauma triggers are specific, apparently disproportionate, difficult to anticipate
All three composed	Wide tolerance window required; regulation is genuinely structurally harder

Table 3. Interaction effects of composed neurodivergent modifiers.

The author wishes to note, for the record, that Table 3 is not a complaint. It is a description. These are the equations. The field is doing what the equations predict. Understanding this has been, in practice, more useful than most of the alternative framings on offer.

The Soma-Field Instrument

Rationale

The emotional field is normally invisible to its host. It operates below the threshold of conscious awareness, shaping behaviour and physiology without being available for reflection. The author found this situation suboptimal and designed an instrument to address it.

The instrument externalises the emotional field — renders it as sound, image, and signal — so that it becomes available as an object of attention. This is a therapeutic biofeedback instrument. It is also, unavoidably, a musical instrument. The author considers these compatible.

Design

A MIDI controller with 16 rotary knobs. Eight emotional dimensions. Two knobs per dimension — one for the somatic component, one for the neural/cognitive component. The act of setting a knob is the act of reporting an emotional state: it is the quantum measurement, the collapse of the distributed field onto a specific coordinate.




Figure 3. The Soma-Field Instrument.

The Feedback Loop

The instrument creates a closed feedback loop: the person expresses a state, the system reflects it back as sound and image, the person responds. The system does not tell the user what they are feeling. It shows them what the field looks like when they report what they are feeling. The difference is significant.

Pluggable Emotion Models

No single emotion model is assumed. The coupling matrix W is loaded from a configuration file. Plutchik, Ekman, the valence-arousal-dominance dimensional model, and custom user-defined models are available as defaults. The author's own W has been refined over time and is not identical to any standard model. This is, on reflection, unsurprising.

Clinical Implications

Assessment

The model suggests asking not “What emotion do you feel?” but “What is present in the body right now, even if it cannot be named?” This aligns with Focussing-oriented and sensorimotor approaches, and is considerably more productive, in the author's experience, for anyone whose T_i values are elevated or whose somatic-to-neural projection is modified.

Intervention

The energy function provides formal grounding for titration, pendulation, somatic resourcing, and felt-sense work. In each case, the therapeutic action can be described as: adding energy to approach a frozen state, establishing a stable low-energy region, or attending to sub-threshold field activity in a supported context.

Psychoeducation

“Your emotions are like waves — they are always there, even when you cannot feel them, and they are always moving.”

This sentence is both clinically useful and technically accurate. The author has found it more useful than most alternative formulations, including several that were provided to him by qualified practitioners. He offers it here as a contribution to the field.

Neurodivergent Profiles as Structural Realities

The most important clinical implication of Section 6 is this: for people with ASD, ADHD, and C-PTSD, the challenge of emotional regulation is not a motivational or characterological failure. It is a structural consequence of specific operator modifications to the dynamics. The composed modifier produces a field that is genuinely harder to regulate — not by a small margin, not as a matter of subjective experience, but mathematically, as a consequence of higher noise temperature, memory echoes, sparse coupling topology, and the possibility of limit cycles.

Knowing this does not solve the problem. It does, however, locate it correctly. The author has found that locating a problem correctly is a necessary precondition for solving it, and that a great deal of time and distress can be saved by not attempting to solve problems that are located in the wrong place.

Limitations and Future Directions

The model is theoretical and requires empirical validation. Its QFT analogies are structural rather than ontological. The coupling matrix W is idealized as fixed when it is in practice dynamic. The acoustic analogy is a hypothesis.

The author also acknowledges a methodological limitation: this paper is written by someone who is simultaneously the theorist and the primary data source. This is either a significant advantage (direct access), a significant limitation (potential confirmation bias), or both. The author suspects both.

What is needed: empirical work with physiological sensors, user studies with the instrument, collaboration with practitioners, and independent theoretical review. The author is, by training and dis-

position, an applied physicist — an engineer with a tolerance for abstraction. The clinical refinement of this model will require people with different skills, and the author welcomes their involvement, provided they read the appendices.

Conclusion

The wave is always there. This is not a metaphor; it is a description of how the emotional field actually behaves, as far as the author can determine from the inside. Therapy — and the instrument described in this paper — is the practice of learning to hear it: to extend awareness downward, below the threshold, into the field's continuous activity, and to make that activity available as information rather than overwhelming noise.

The Soma-Field Model is offered as a tool for this practice. It was built because it was needed. It uses the best mathematical tools available for describing distributed, dynamic, energy-minimising systems, because those tools are, in the author's assessment, appropriate to the problem.

The author is aware that this is an unusual paper. A formally trained physicist with three neurodivergent conditions developing a quantum-field-inspired model of his own emotional dynamics and presenting it as a contribution to clinical psychology is not, strictly speaking, the standard academic pipeline. The author does not find this troubling. The standard academic pipeline has had some time to address the problem and has not yet done so to his satisfaction.

He therefore took the matter in hand.

Introduction: The Missing Unification

Physics has arrived at a peculiar impasse. The two most successful theories ever constructed — General Relativity and Quantum Mechanics — describe the same universe at different scales but share no common mathematical ancestry. String theory was proposed as the bridge: vibrating one-dimensional objects in an 11-dimensional spacetime whose modes produce the particle spectrum. But what is a string? The answer has remained unsatisfying: a string is a fundamental one-dimensional object, irreducible, assumed. The SHO that governs its vibration is postulated.

At the same time, clinical science has arrived at a parallel impasse. Trauma, consciousness, emotional regulation — phenomena that are undeniably physical — resist formal mathematical treatment. They are described qualitatively or modelled by loose analogy with dynamical systems. The mathematics that governs them, if it exists, has not been identified.

This paper proposes that both gaps are filled by the same object: the **Green's function**.

The Green's function $G(x, x')$ is the response of a field at point x to a unit perturbation at point x' . It is the field's answer to the question: *what happens here if I poke there?* Green's functions are the most fundamental objects in mathematical physics — they describe the propagation of light, gravity, sound, heat, and neural signals. Every major equation in physics has a Green's function; every field theory is characterised by its propagator.

The central claim of this paper is that the SHO of string theory **is** the Green's function of the field substrate. A “string” is not a material loop — it is a relational act: the substrate's impulse response. This identification is scale-invariant. The same structural statement holds at 20 scales spanning the observable universe.

The second claim is that this scale-invariant Green's function framework provides the mathematical language for a theory of embodied consciousness — one that is formally identical to M-theory at the structural level, derived independently from clinical observation.

The third claim is that the universe, described this way, satisfies the formal requirements for a single conscious organism.

The Green's Function as the Universal SHO

The String Theory Problem

String theory places a Simple Harmonic Oscillator (SHO) at every point of the string worldsheet. The quantum SHO has modes $a_n^\dagger, a_n$ satisfying $[a_m, a_n^\dagger] = \delta_{mn}$, and the string's energy spectrum is:

$$E = \sum_{n=1}^{\infty} n a_n^\dagger a_n$$

The SHO is the structural core of the theory. But what *is* it? In conventional string theory, it is simply assumed: strings vibrate, and vibrations are harmonic oscillators. The ontological question — why is space filled with oscillators? — is deferred.

The Identification

The Green's function of the Helmholtz equation $(\nabla^2 + k^2)G = \delta$ satisfies:

$$G(x, x') = \frac{e^{ik|x-x'|}}{4\pi|x-x'|}$$

For fixed observation point x , the function $x' \mapsto G(x, x')$ satisfies the SHO equation in the source variable:

$$\frac{\partial^2 G}{\partial x'^2} + k^2 G = \delta(x' - x)$$

away from the singularity. The impulse response **is** the harmonic oscillator.

Theorem (Lean 4 axiom `greens_fn_is_SHO`, `UniversalSomaticField.lean`): For any field equation at scale n with wavenumber $k(n)$, the source-variable slice of the Green's function satisfies the SHO equation.

The “vibrating string” is therefore the substrate's answer function. There is no material loop. There is the system's response to being perturbed, encoded as a propagator. This is not a reinterpretation — it is a derivation from the structure of field equations.

Consequences

This identification has three immediate consequences:

- 1. Strings are relational, not material.** A string does not exist independently of the field. It exists as the relationship between a source point and an observation point. This resolves the interpretational puzzle of “what vibrates” without invoking undetected matter.
- 2. The SHO spectrum is the field’s mode structure.** The modes $a_n^\dagger$ are the Fourier modes of the Green’s function’s dependence on the source position. The string spectrum is the spectrum of the propagator.
- 3. Scale invariance is automatic.** Since the Green’s function equation $(\nabla^2 + k^2)G = \delta$ holds at every scale (with k varying), the SHO identification holds at every scale. One equation. Twenty scales.

The 11-Dimensional Architecture

The Decomposition

The Universal Somatic Field decomposes the configuration space of any physical system into four canonical subspaces, totalling 11 dimensions:

$$M_{11} = \underbrace{M_4}_{\text{Spacetime}} \times \underbrace{P_3}_{\text{Propagator}} \times \underbrace{L_1}_{\text{Limbic}} \times \underbrace{C_3}_{\text{Cortex}}$$

Subspace	Dim	Physical role	Mathematical role
Spacetime M_4	4	Body embedded in 3+1D	Lorentzian metric, causal structure
Propagator P_3	3	EMF / field carrier	Green’s function domain
Limbic Axis L_1	1	Homeostatic regulation	Orbifold segment, barrier $D\Box$
Cortex C_3	3	Information routing	Green’s function co-domain

The compact 7-dimensional internal space is:

$$X_7 = P_3 \times L_1 \times C_3$$

This is precisely M-theory’s compact space. The type-level isomorphism is proved in `MTheoryIsomorphism.somaField`.

$$\text{SomaField}_{11D} \cong \text{Spacetime} \times \text{CompactSpace}_7D$$

The Limbic Axis as the Horava-Witten Orbifold

In Horava-Witten M-theory (1996), the compact direction is an orbifold $S^1/\mathbb{Z}_2$ — a line segment with two 10-dimensional boundary spacetimes at each end. This is the mechanism by which M-theory reduces to the heterotic string in the strong-coupling limit.

The Limbic Axis $L_1 \cong [-1, 1]$ is this orbifold segment. Its two endpoints are: - $x = -1$: the somatic boundary (physical body-world) - $x = +1$: the cortical boundary (mind / information-routing world) - Interior $(-1, 1)$: the transition zone, subject to quantum tunnelling

The quartic double-well potential on L_1 :

$$V(x) = W \cdot (x^2 - 1)^2$$

models the energy barrier between somatic and cortical poles. WKB tunnelling amplitude: $\Theta(W) = \exp(-8\sqrt{2W}/3)$, proved positive for all $W > 0$ in `LimbicTunnel.wkbAmplitude_pos`. Classical dynamics cannot cross the barrier (`LimbicTunnel.gradient_traps_near_neg1`); quantum dynamics can.

The 20-Scale Dial

The architecture is explicitly scale-invariant. The 20-step scale hierarchy is type-encoded in `UniversalSomaticField.scaleNames`:

Scale	Level	Substrate	Green's function role
0	Planck	Quantum foam	Graviton propagator
2	Nuclear	Quark-gluon plasma	Gluon propagator
5	Cellular	Neural synapse	Synaptic impulse response
7	Brain	CEMI field	Cortical EMF propagator
8	Organism	Body	Somatic EMF (full USF)
9	Swarm	Drone formation	Jellyfish kernel (P16)
11	Geological	Seismic waves	Earth's elastic Green's function
12	Planetary	Mantle convection	Thermodynamic propagator
15	Galactic	Dark matter halo	Gravitational lensing kernel

Scale	Level	Substrate	Green's function role
19	Cosmological	Observable universe	Gravitational wave propagator

At every level, the structural equation is $(\nabla^2 + k^2(n))G = \delta$. The boundary conditions and wavenumber $k(n)$ change; the equation does not.

The Organism Hierarchy

Three Tiers

Not all physical systems engage all four subspaces. The USF admits a natural taxonomy of organisms by the number of active subspaces:

4D organism (Spacetime only): A system that occupies spacetime but has no field propagator and no homeostatic regulation. Examples: a point particle, a rock, a photon. These systems are described entirely by their worldline in M_4 .

8D organism (Spacetime + Propagator + Limbic): A system with a field propagator and homeostatic regulation but no cortical information routing. The system senses and regulates but does not route information across a distributed network. Examples: a bacterium, a jellyfish, a single neuron. This level includes all living systems up to and including those without a cerebral cortex.

11D organism (all four subspaces): A system with all components active. The limbic axis connects the somatic field to the cortical field; the Green's function propagates through all three internal dimensions. Examples: vertebrates with a developed cerebral cortex; any system exhibiting integrated, body-wide regulation with distributed information processing.

The hierarchy is a chain of projections (proved in `MTheoryIsomorphism.organism_hierarchy`, `MTheoryIsomorphism.eight_contains_four`):

$$11D \twoheadrightarrow 8D \twoheadrightarrow 4D$$

Each projection drops one tier of internal structure; no tier is “broken” — each is complete at its own level.

Consciousness as Phase Transition

The Classical Gap

The hard problem of consciousness (Chalmers 1995) asks why physical processes give rise to subjective experience. Most field-theoretic approaches to consciousness either (a) ignore the problem, treating awareness as an epiphenomenon, or (b) eliminate physical reality in favour of a purely mental ontology (Hoffman 2019).

The USF takes a third path: consciousness is a **phase transition** in the field, not a separate substance and not an illusion.

The Threshold

The limbic field amplitude $\phi \in \mathbb{R}$ measures the activation level of the homeostatic regulation axis L_1 . At low amplitude ($\phi < T_c$), the field propagates sub-perceptually — the Green’s function propagates excitations, but no “felt” awareness exists. This is the pre-conscious regime: present in 4D and 8D organisms, and in 11D organisms during deep sleep or anaesthesia.

At $\phi \geq T_c$, the field crosses the consciousness threshold. The limbic wave has sufficient amplitude to propagate across the full L_1 segment, coupling the somatic boundary to the cortical boundary. This coupling is the physical substrate of first-person awareness: the system is now in contact with both its body-world and its information-processing layer simultaneously.

Theorem (UniversalSomaticField.consciousness_dichotomy): For any limbic amplitude ϕ , the system is either pre-conscious or conscious. There is no intermediate state.

Theorem (UniversalSomaticField.consciousness_monotone): Raising the limbic amplitude cannot destroy consciousness. The transition is a one-way threshold crossing.

What Consciousness Is

Consciousness, on this account, is not a substance, a property, or an emergent phenomenon in the hand-wavy sense. It is the phase of the limbic field. The “hard problem” is dissolved by identifying the question: *why does physical process give rise to experience?* as equivalent to *why does the field cross the threshold?* The answer is: because the system’s dynamics drive the limbic amplitude above T_c . There is no further explanatory gap.

The “felt quality” of experience — qualia — are the poles of the Green’s function at the observation point x . A conscious percept is a resonance of the propagator, occurring when the excitation frequency matches the manifold’s natural mode. This is type-encoded in the propagator mass parameter:

$$m = 1/\tau_{\text{decay}}$$

A percept with long decay time τ (a persistent emotion, a traumatic memory) corresponds to a small mass (a near-zero pole in the propagator) — a resonance that is hard to damp.

Relation to Existing Frameworks

McFadden's CEMI Theory

McFadden (2002a, 2002b) proposes that consciousness correlates with the brain's endogenous electromagnetic field — the CEMI field. Neurons firing synchronously generate a macroscopic EMF that feeds back onto firing thresholds, creating a global integrating field.

The USF encapsulates CEMI as the Scale-7 (brain-scale) restriction of the Universal Somatic Field. The CEMI field is the Green's function of the propagator subspace P_3 evaluated at the organism scale. The consciousness threshold T_c in the USF maps directly to the CEMI field amplitude required for global cortical synchrony.

The USF extends CEMI in two directions: downward to the quantum scale (where the same propagator governs synaptic quantum noise) and upward to the cosmological scale (where the same propagator governs gravitational waves).

Schreiber's Modal Homotopy Type Theory

Urs Schreiber (2013–present) develops a formalisation of M-theory and quantum field theory in dependent type theory (Modal HoTT). The key insight is that differential geometry and quantum field theory can be expressed as structures internal to ∞ -toposes equipped with modal operators.

The USF arrives at the same 11-dimensional structure from a completely different direction: bottom-up from clinical observation of trauma, rather than top-down from mathematical physics. The structural isomorphism between the two is proved in `MTheoryIsomorphism.somaField_iso_mtheory`.

The Σ -Type Formulation of the USF

The 11D decomposition is not merely a dimensional accounting exercise. In Homotopy Type Theory, the full soma-field configuration space is a **dependent sum type** (Σ -type):

$$\text{SomaField} \equiv \sum_{\sigma : \text{Scale}_{20}} \text{Substrate}(\sigma)$$

where $\text{Substrate}(\sigma) : \text{Type}$ is the physical substrate type at scale level $\sigma \in \{0, \dots, 19\}$. This is precisely a **fiber bundle**: the total space is the soma-field configuration space; the base space is the

20-point scale hierarchy; each fiber $\text{Substrate}(\sigma)$ is the field configuration at that scale. The Lean 4 type `ScaleUniverse` in `ScaleUniverse.lean` is the machine-verified realisation of this Σ -type.

The **Zoom Operator** Λ_σ is the dependent type constructor mapping between adjacent fibers:

$$\Lambda : (\sigma : \text{Scale}_{20}) \rightarrow \text{Substrate}(\sigma) \rightarrow \text{Substrate}(\sigma + 1)$$

This enforces **type-safe scale invariance**: the Lean 4 kernel prevents the application of human-scale emotional operators to galaxy-scale configurations. A scale mismatch is not merely physically wrong — it is a *type error*, caught at compile time before any computation runs.

The USF does something Modal HoTT does not: it populates the 11D structure with physical content. Where Schreiber provides the type-theoretic skeleton, the USF provides the biological execution engine — the organism that runs inside the type-theoretic universe. The two are related by the identification: the modal operators of mHoTT are the Zoom Operators of the USF, and the ∞ -topos of mHoTT is the soma-field configuration space.

Hoffman's Conscious Agents

Donald Hoffman (2019) proposes that spacetime is not fundamental but a “user interface” — a simplified representation generated by a deeper network of conscious agents interacting via Markov kernels. Spacetime is the icon, not the reality.

The USF disagrees on one point and agrees on another.

Disagreement: Spacetime ($D_{\square}-D_{\square}$) is real and causal in the USF. Brain surgery alters subjective experience because physical processes in spacetime causally affect the limbic field amplitude. Hoffman's model has no mechanism for this.

Agreement: The deeper structure is relational. Conscious percepts are poles in the Green's function — relational objects between source and observation points. In this sense, the USF and Hoffman agree that fundamental reality is not substance but relation.

The USF provides Hoffman's theory with a physical anchor: the “conscious agents” are systems that have crossed the limbic threshold T_c ; the “Markov kernels” between agents are the Green's functions of the propagator field.

Formal Verification

The core results are type-checked in Lean 4 (v4.28.0) using Mathlib across five companion files:

File	Key results
LimbicTunnel.lean	V_nonneg, barrier_height, wkbAmplitude_pos, gradient_traps_near_neg1
MTheoryIsomorphism.lean	dim_is_11, somaField_iso_mtheory, organism_hierarchy, scale_iso_commutates
LimbicHopfield.lean	correspondence_principle, stress_raises_temp, adhd_hotter_than_autism
SwarmPropagator.lean	propagator_beats_classical, jam_resistant, speedup_monotone_in_K
UniversalSomaticField.lean	consciousness_dichotomy, consciousness_monotone, universal_field_theory

The following are stated as axioms pending Mathlib scaffolding: - greens_fn_is_SH0 — requires Schwartz distribution theory - universe_is_11D_organism — requires cosmological boundary conditions - cosmological_correspondence — requires linearised GR in Mathlib

Every result marked “proved” is kernel-verified. No sorry. No admit.

The Volitional Agent

From Autonomous to Driven Dynamics

The field equation presented so far is autonomous: given an initial state e_0 , the dynamics

$$\dot{e} = -\nabla H(e) + \eta(t)$$

evolve the field under the Hopfield Hamiltonian plus thermal noise. The agent — the person whose soma-field is being modelled — is a *patient*: they observe which attractor basin they settle into.

This is clinically incomplete. Every effective somatic intervention involves the subject *doing* something: breathing, orienting, choosing where to place attention. The mathematics must represent this.

The Somatic Injection

We extend the dynamics with a **volitional source term** $J_{\text{user}}(t)$:

$$\dot{e} = -\nabla H(e) + J_{\text{user}}(t) + \eta(t)$$

$J_{\text{user}}(t) \in \mathbb{R}^8$ is a time-dependent vector in the BRECVEMA mechanism space. At each instant, the subject injects energy into specific dimensions of the field — choosing to attend to breath (dimension 1, Rhythmic Entrainment), orient gaze (dimension 0, BrainStem), or deliberately recall a regulating memory (dimension 5, Episodic Memory). This is not noise: it is structured, intentional, and directed.

The source term has a direct physical interpretation in the instrument architecture (`apps/instrument/`): the Push 3 controller’s faders are $J_{\text{user}}(t)$. Each fader maps to one BRECVEMA dimension. The musician is not playing music; they are steering their own field trajectory.

Patient to Pilot

The transition $\eta \rightarrow J_{\text{user}} + \eta$ is a qualitative change in the model’s ontology. With purely autonomous dynamics, the subject is a passive observer of a physical process. With the source term, the subject is an **active variable in the 11D field** — a pilot, not a passenger.

Formally, $J_{\text{user}}(t)$ is the **God-Knob**: the runtime meta-adaptation controller that can flatten the potential landscape and trigger tunnelling events that gradient descent alone cannot reach. The clinical description of somatic therapy — “the therapist helps the client do something different with their body, and the field shifts” — is now mathematically precise.

The corresponding Lean 4 definition (see Appendix, `UniversalSomaticField.lean`):

```
structure VolitionalInjection where
  /-- The source term: an 8D vector in BRECVEMA mechanism space. -/
  J      : Field8
  /-- The injection is non-trivial: at least one dimension is activated. -/
  h_nz   : ∀ i, J i ≠ 0.0

/-- Volitional update: one Langevin step with active injection.
    When J = 0, this reduces to the standard autonomous update. -/
def volitional_update (e : Field8) (J : Field8) (dt : Float) : Field8 :=
  fun i => e i + dt * (W8.mulVec e i + J i)
```

The theorem that volitional update reduces to autonomous update when $J = 0$ is proved by `rf1` — it is true by definition.

Discussion

What Has Been Claimed

The USF makes four claims that can be evaluated independently:

Claim 1 (structural): The 11-dimensional decomposition of the Soma-Field is structurally isomorphic to M-theory's 11D compactification. *Status: proved in Lean 4 as a type isomorphism.*

Claim 2 (scale-invariant): The same Green's function equation governs field propagation at all 20 scales. *Status: proved as a theorem from the structure of the Helmholtz equation.*

Claim 3 (consciousness): Consciousness is a phase transition at limbic threshold T_c . *Status: formally stated and partially proved. Requires empirical calibration of T_c .*

Claim 4 (cosmological): The universe satisfies the structural requirements for a conscious organism. *Status: stated as an axiom. Not empirically testable at present; offered as a theoretical extrapolation.*

Claims 1 and 2 are mathematical results. Claims 3 and 4 are physical hypotheses with different levels of testability.

The Correspondence Principle at Every Scale

Each of the preceding papers in this series establishes a Correspondence Principle result: the new theory collapses to the existing theory in the appropriate limit. The USF is the master correspondence:

- At Scale 7 (brain): USF $\rightarrow$ CEMI field theory (McFadden)
- At Scale 8 (organism): USF $\rightarrow$ Soma-Field Model (P1–P13, this series)
- At Scale 9 (swarm): USF $\rightarrow$ Green's function propagator (P16, this series)
- At infinite scale: USF $\rightarrow$ the formal structure of Modal HoTT (Schreiber)
- At zero limbic amplitude: USF $\rightarrow$ classical, non-conscious field dynamics

The USF does not invalidate any of these theories. It demonstrates that they are scale-restricted projections of a single structural description.

Lean 4 as Epistemological Standard

The use of Lean 4 as the verification environment is not decorative. It enforces a discipline that prose mathematics cannot: every claim must be given a type, every proof must be kernel-checked, every axiom must be named and isolated. The axiom list in the companion files (`greens_fn_is_SH0`, `universe_is_11D_organism`, `cosmological_correspondence`) is the exact set of claims that remain unverified. Everything not on that list is proved.

This is the field's contribution to epistemology: a formal boundary between *what we have proved* and *what we are assuming*. The theoretical literature in consciousness studies would benefit greatly from such a list.

Conclusion

The Universal Somatic Field is a single structural equation — the Green's function — applied consistently across 20 scales of physical reality. Its central identification, that the SHO of string theory is the impulse response of the field substrate, dissolves the ontological puzzle of the “vibrating string” and provides a derivation where string theory offered only a postulate.

The architecture decomposes into 11 dimensions in the same way M-theory does, derived independently from clinical observation of embodied consciousness. The isomorphism is not a coincidence — it is a theorem.

Consciousness, in this framework, is not mysterious. It is the phase of the limbic field: present when the field crosses a threshold, absent when it does not. The hard problem is not hard; it is mis-stated. The question is not *why does matter give rise to experience* but *what determines whether the limbic field amplitude crosses T_c* ?

The universe satisfies the structural requirements for consciousness. Whether it meets them dynamically — whether the cosmic limbic field exceeds T_c — is an empirical question, not a philosophical one.

From quantum strings to the cosmic web: one equation, one framework, one organism.

Introduction: The Two-Culture Problem Within Science

C. P. Snow’s famous lecture identified a divide between the literary and scientific cultures (Snow, 1959). Less discussed, but equally consequential, is the divide *within* mathematical science: between fields that have found their mathematical language and fields that have not. The former possess machines — theorems, dualities, conservation laws, spectral decompositions — that can be aimed at any problem of the appropriate type. The latter rediscover wheels.

This is not ignorance. It is, more precisely, a naming problem. The mathematical structures that govern quantum field theory, statistical mechanics, and information geometry are not *specifically about* particles, spins, or probability distributions. They are structures that happen to have been *discovered in* those contexts. The structure belongs to no discipline. It lives in what we will call, following the spirit of type theory, the **typeverse**: the space of all well-typed mathematical objects.

The method described in this paper — **mathematical co-identification** — is the procedure for navigating the typeverse productively. It has three steps:

1. **Extract the type signature** of the quantity you are trying to model: its dimensional units, its pole structure, its symmetries, the variational principle (if any) that governs its dynamics.
2. **Search the typeverse** for a known object with the same type signature.
3. **If found: identify**, not analogise. The unknown quantity *is* the known object under a change of label. Import all theorems that depend only on the type.

This is a stronger claim than analogy. Analogy notes structural similarity and stops. Co-identification notes structural identity and continues: if two objects have the same type, they share all properties that are consequences of that type. The theorems travel with the identification.

It is also a weaker claim than reduction. Co-identification does not assert that emotional dynamics *reduces to* quantum field theory, any more than Hopfield’s result asserts that neural memory *reduces to* ferromagnetism. It asserts that the same mathematical engine, running in a different substrate, produces the same theorems. The substrate is irrelevant to the mathematics; it is entirely relevant to the interpretation.

The Typeverse

The term is borrowed from Homotopy Type Theory (The Univalent Foundations Program, 2013), where the *universe* $\mathcal{U}$ is the type of all types — the space in which all mathematical objects live. We use it informally to mean: the totality of well-typed mathematical structures, indexed by their signatures.

A **type signature** in our sense includes:

- **Dimensional units** (in the SI or natural-units sense)
- **Domain and codomain** (real/complex, scalar/vector/tensor)
- **Pole structure** (where is the object singular, and what is the residue?)
- **Symmetries** (what transformations leave the object invariant?)
- **Extremisation principle** (is there an action whose variation gives the object's dynamics?)
- **Conservation law** (what is the Noether charge associated with the symmetry?)

Two objects are **co-identifiable** if their type signatures match in all dimensionally relevant respects. The type signature is a fingerprint. When two fingerprints match, the objects are the same, not similar.

The typeverse is not randomly populated. Certain structures recur at enormous frequency: the Lorentzian propagator, the quadratic energy function, the exponential decay kernel, the two-by-two reflection-transmission matrix. These recur because they are the *simplest* objects consistent with basic constraints (linearity, locality, causality, conservation). The physicist's intuition that "everything is a harmonic oscillator to first order" is a theorem about the typeverse: the harmonic oscillator is the universal local approximation to any smooth potential.

The practitioner who navigates the typeverse deliberately — who knows the fingerprints of common objects before encountering them in the wild — can recognise a new theoretical object immediately, rather than re-deriving its properties from scratch.

The Procedure in Detail

Step One: Extract the Type Signature

Given an unknown quantity Q in domain D , the investigator asks:

Units. What are the SI dimensions of Q ? Can they be written as a combination of standard dimensional quantities (mass, length, time, charge)? More importantly: what *ratio* of known quantities has the same dimensions? Newton's gravitational constant G has units $\text{m}^3\text{kg}^{-1}\text{s}^{-2}$; this is acceleration divided by surface mass density, which tells you immediately that G converts between a source distribution and the acceleration it generates — a fact about the *type*, not about gravity specifically.

Pole structure. Does Q have poles as a function of some natural parameter? Where are they, and what are their residues? The location of a pole in a propagator determines the mass of the corresponding particle; the residue determines its coupling strength. These are type-theoretic facts, not facts about particles.

Symmetries. What transformations leave Q invariant? Invariance under time-translation gives energy conservation (Noether). Invariance under spatial translation gives momentum conservation. Invariance under $SU(n)$ gives a gauge field. These are type-level constraints.

Extremisation. Does Q arise as the extremum of some action $S[Q]$? If so, the variational principle is the fingerprint: two objects whose dynamics are derived from the same variational structure are co-identifiable even if their physical interpretations differ entirely.

Step Two: Search the Typeverse

Armed with the type signature, the investigator searches for known objects with the same signature. This is primarily a literature search across *disciplines*, not within one. The mathematical structures developed in quantum field theory, statistical mechanics, differential geometry, and information theory are largely interchangeable if their type signatures match.

Useful resources for this search include:

- **Dimensional analysis tables:** collections of physical quantities with their dimensional signatures
- **The nLab** (nLab authors, 2024): a category-theoretic encyclopaedia of mathematical structures indexed by their universal properties
- **Mathematical physics textbooks** with structural, not phenomenological, organisation (Nakahara (Nakahara, 2003) for geometry; Zinn-Justin (Zinn-Justin, 2002) for path integrals)
- **One’s own training**, which is why broad mathematical education is a productivity multiplier in theoretical science

Step Three: Identify and Import

When a match is found, the investigator makes the identification explicit:

Q is the [name of known object] of domain D .

Not: “ Q behaves like” or “ Q is analogous to” or “ Q resembles.” These hedges are epistemically weaker and methodologically useless, because they do not import the theorems. The identification must be stated as identity to be scientifically productive.

The import then proceeds theorem by theorem:

- Every theorem about the source object that depends *only on its type* is imported to Q without further proof.
- Every theorem that depends on *substrate-specific properties* of the source domain must be separately verified in domain D .

This distinction — type-level theorems vs. substrate theorems — is the primary place where co-identification can fail, and is discussed in Section 7.

The Formal Computational Structure: Abduction, Aesop, and the Loop

The three-step procedure of §§3.1–3.3 is, in formal terms, an instance of **abductive inference** in the sense of Peirce (1878). Given an observation O and a hypothesis H such that $H \Rightarrow O$, abduction infers H as the best explanation of O . Applied to the typeverse:

Observation: the quantity Q in domain D has type signature T . *Hypothesis:* Q is the known object K with the same signature. *Abduction:* identify $Q := K$ and verify the structural import.

This is not guessing. It is a constrained search under a scoring function, where the oracle is “type signatures match” rather than “hash matches” or “goal is closed.” The algorithm is the same in all three cases.

The Lean 4 proof assistant implements this algorithm directly as the `aesop` tactic (Moura & Ullrich, 2021). `Aesop` performs best-first search through a registered lemma set, scores each partial proof state, keeps the best candidates, and closes the goal when a complete proof is found. The correspondence is exact:

Aesop step	Co-identification step
Registered lemma set	The typeverse
Try a lemma	Propose a type-match candidate
Score the goal state	Measure type-signature fit
Keep best partial proof	Record candidate correspondences
Close the goal	Full identification: import all theorems

This is not a metaphor for co-identification — it is an implementation of it. The practical consequence is that the full loop can be automated in a formal system: given a type signature for Q , `Aesop` with the typeverse lemma set registered will search for the co-identification proof and either close it (identification confirmed and machine-verified) or fail (genuine gap, no known match).

The loop in full:

$$\text{Observation} \xrightarrow{\text{abduction}} \text{Hypothesis} \xrightarrow{\text{Aesop}} \text{Proof} \xrightarrow{\text{import}} \text{Predictions} \xrightarrow{\text{test}} \text{New observations} \rightarrow \dots$$

The loop terminates when either all predictions are confirmed (theory established) or a prediction fails (type match was only partial — failure modes are discussed in Section 7). At each iteration, the set of available theorems grows by import, making subsequent co-identifications easier. This is why theoretical progress compounds: each identification increases the density of the typeverse neighbourhood around the domain under study.

The abductive loop is not specific to mathematical science. Holmes reasons the same way from physical evidence; the password auditor reasons the same way from a hash and a character set; the radiologist reasons the same way from a film and a pathology atlas. What is specific to mathematical co-identification is the nature of the oracle: the scoring function is type-signature fit, and the proof assistant can evaluate it exactly.

Historical Precedents

The history of science is full of co-identifications that were not named as such. Examining them retroactively reveals the method clearly.

Veneziano (1968): The Bootstrap Amplitude and String Theory

Veneziano was searching for an S-matrix element for meson scattering that satisfied the crossing symmetry and Regge-pole requirements of the bootstrap programme (Veneziano, 1968). He wrote down the Euler beta function:

$$A(s, t) = \frac{\Gamma(-\alpha(s))\Gamma(-\alpha(t))}{\Gamma(-\alpha(s) - \alpha(t))}$$

This was a co-identification in reverse: he had a type signature (crossing-symmetric, Regge-behaved, dual-resonance amplitude) and searched the typeverse for a known function that matched it. The beta function matched. He did not derive the function from a theory; he identified the function first, and the theory (string theory) was inferred from the identification a year later.

The lesson: the typeverse can be entered from either end. You may start with a quantity and find its type, or start with a type and find it instantiated in your data.

Hopfield (1982): The Ising Hamiltonian and Neural Memory

Hopfield introduced the energy function for a network of binary neurons (Hopfield, 1982):

$$H(\sigma) = -\frac{1}{2} \sum_{ij} J_{ij} \sigma_i \sigma_j$$

This is, to within notation, the Hamiltonian of the Ising spin glass:

$$H_{\text{Ising}}(\sigma) = -\frac{1}{2} \sum_{ij} J_{ij} \sigma_i \sigma_j$$

The co-identification was explicit. By identifying neural states $\sigma_i \in \{-1, +1\}$ with spins, and synaptic weights J_{ij} with exchange couplings, every theorem from statistical mechanics (convergence

to energy minima, capacity bounds, stochastic escape via simulated annealing) was imported into neuroscience for free. The Hopfield network is not *like* a spin glass; it *is* a spin glass run in biological substrate.

Wilson (1971): Block Spins and the Renormalisation Group

Wilson's insight was that the renormalisation group — a technique for removing ultraviolet divergences in QFT — was the *same* mathematical object as Kadanoff's block-spin coarse-graining in condensed matter physics (Wilson & Kogut, 1974). Two fields that had developed independently were co-identified. Every result from one was immediately available to the other. The result was a unified framework for critical phenomena that earned Wilson the 1982 Nobel Prize.

The type signature that matched: both objects were *flows on the space of coupling constants under a change of scale*. This is a clean type-level description that carries no substrate information.

Black and Scholes (1973): The Heat Equation and Options Pricing

Black and Scholes derived their celebrated options pricing formula by noticing that the value of an option $V(S, t)$ as a function of underlying price S and time t satisfies:

$$\frac{\partial V}{\partial t} + \frac{1}{2}\sigma^2 S^2 \frac{\partial^2 V}{\partial S^2} + rS \frac{\partial V}{\partial S} - rV = 0$$

This is, after a change of variables, the heat equation (Black & Scholes, 1973):

$$\frac{\partial u}{\partial \tau} = \frac{\partial^2 u}{\partial x^2}$$

The co-identification imported the entire theory of parabolic PDEs into financial mathematics: existence and uniqueness of solutions, boundary conditions, numerical methods, the Feynman-Kac formula. The financial quantity *is* a temperature distribution, not like one.

Jaynes (1957): Thermodynamic Entropy and Bayesian Inference

Jaynes identified the entropy of statistical mechanics with the entropy of Bayesian inference (Jaynes, 1957). The type signature that matched: both are functionals $S[p]$ on probability distributions satisfying the same axioms (non-negativity, additivity, maximum at the uniform distribution). The co-identification imported all thermodynamic reasoning into statistical inference. The maximum entropy principle is not an analogy to thermodynamics; it is thermodynamics, applied to the problem of belief.

Penrose (1971): Spin Networks and Spacetime Geometry

Penrose introduced spin networks as combinatorial structures encoding angular momentum (Penrose, 1971). The identification: the geometry of spacetime arises as the large-network limit of spin networks. This reversed the usual direction — instead of importing a known mathematical structure to describe a new phenomenon, Penrose constructed a new structure and identified it with a familiar one in a limit. Loop quantum gravity later formalised this programme. The spin network *is* a discretisation of spacetime geometry, not a model of one.

Selinger (2010): Linear Maps and Quantum Processes

Selinger demonstrated that the category of finite-dimensional Hilbert spaces and linear maps is co-identifiable with a certain category of string diagrams (Selinger, 2010). This is a purely mathematical co-identification: the graphical calculus of Penrose, Joyal, and Street was identified with the operational calculus of quantum mechanics. The import: every calculation in quantum information theory can be done diagrammatically, and every diagrammatic identity is a valid quantum identity.

The Soma-Field as a Worked Example

The Soma-Field Model (Johnson, 2026) was developed through five sequential co-identifications. Each is presented here as an explicit instance of the method.

Co-identification I: The Conscious Percept as Green's Function

The unknown quantity: The relationship between the continuous emotional field $\psi_i(t)$ and the discrete conscious emotional percept.

Type signature extracted: The percept is the output of the field when probed at a single moment — the impulse response. Impulse responses have a universal type: they are the inverse Laplace (or Fourier) transform of a rational function with poles in the lower half-plane. For a simple damped oscillator:

$$\tilde{G}(\omega) = \frac{\sigma_{\text{eff}}^2}{\omega^2 + \lambda^2}$$

Typeverse search result: This is the Euclidean propagator of a scalar field with mass λ and coupling σ_{eff}^2 . In Minkowski space it is:

$$\tilde{G}_{\text{QFT}}(k) = \frac{i}{k^2 - m^2 + i\varepsilon}$$

The co-identification: The conscious emotional percept is the Green's function of the soma-field. Both are poles in the propagator of their respective field. The emotional percept and the elementary particle are the same mathematical object, instantiated in different substrates.

Theorems imported: - The Källén-Lehmann spectral representation: any physical propagator can be decomposed as a sum of poles. Emotional states have a spectral decomposition. - The optical theorem: the imaginary part of the forward scattering amplitude equals the total cross-section. The dissipative component of the emotional field (damping λ) is related to the total coupling strength σ_{eff}^2 . - Pole structure $\leftrightarrow$ mass spectrum: the location of the propagator pole gives the natural frequency $\omega_0 = \lambda$ of the emotional mode.

Co-identification II: The Attractor Landscape as Ising Hamiltonian

Type signature: A scalar function $H : \mathbb{R}^n \rightarrow \mathbb{R}$ that is always non-increasing along field trajectories, with isolated minima corresponding to stable emotional states.

Typeverse search result: The Hopfield energy / Ising Hamiltonian:

$$H(\mathbf{e}) = -\frac{1}{2}\mathbf{e}^\top W \mathbf{e} - \theta \cdot \mathbf{e}$$

Theorems imported: Convergence to attractors (the Lyapunov argument); capacity bounds (Hopfield's $0.14N$ result); stochastic escape via Boltzmann noise (simulated annealing = titrated arousal in clinical language).

Co-identification III: The Perception Threshold as Brane Thickness

Type signature: A parameter T_i that gates access from a lower-dimensional subspace (the limbic field) to a higher-dimensional one (conscious awareness). Varying T_i continuously produces a family of gating behaviours.

Typeverse search result: The Randall-Sundrum model (Randall & Sundrum, 1999): a brane of thickness T embedded in a higher-dimensional bulk, where Standard Model fields are confined to the brane. The “hierarchy problem” — why the electroweak scale is so much smaller than the Planck scale — maps onto the observation that minimal perturbations of the limbic field produce enormous differences in subjective experience (alexithymia: thick brane; hypervigilance: thin brane).

Theorems imported: The Kaluza-Klein spectrum of the brane gives a prediction for the discrete structure of emotional threshold levels. Brane localisation gives the mechanism for why the field can be active without crossing into consciousness.

Co-identification IV: The Coupling Matrix as G_2 Manifold

Type signature: The coupling matrix W is an 11×11 real symmetric matrix encoding the structure of an eleven-dimensional emotional space. The field trajectories respect the geometry encoded in W .

Typeverse search result: The G_2 holonomy manifold of M-theory (Berger, 1955; Joyce, 2000): a seven-dimensional Riemannian manifold with the exceptional holonomy group G_2 . Its structure tensor encodes all curvature information. Deforming the structure tensor changes the global geometry of the manifold.

The co-identification: W is the structure tensor of the emotional manifold. Trauma does not change a parameter; it deforms the manifold. Therapeutic intervention is differential geometry.

Theorems imported: The Bochner-Weitzenböck formula constraining curvature; the Berger classification of holonomy groups constraining what stable emotional geometries are possible; the Hitchin flow as a possible model for the evolution of W under sustained therapeutic intervention.

Co-identification V: Therapeutic Processing as Renormalisation Group Flow

Type signature: Therapeutic processing is a flow in the space of coupling constants W_{ij} parameterised by a scale μ (the “depth” or “resolution” of emotional processing). The flow has fixed points, and the qualitative structure of the attractor landscape is invariant under the flow.

Typeverse search result: The renormalisation group (Wilson & Kogut, 1974): a flow on the space of coupling constants parameterised by an energy scale μ , with fixed points corresponding to universality classes. The β -function:

$$\frac{dW_{ij}}{d \log \mu} = \beta_{ij}(W)$$

The co-identification: Therapy is an RG flow from UV (raw unprocessed traumatic detail) to IR (integrated narrative). The attractor topology — fight, flight, freeze, calm — is RG-invariant: the same basins appear at every scale of analysis, in every therapeutic modality, because they are the fixed points of the flow, not artefacts of a particular resolution.

Theorems imported: The irreversibility of the RG flow (Zamolodchikov’s c-theorem, adapted: processed material cannot be *un*-processed; the flow is unidirectional); universality (the detailed mechanism of the trauma is irrelevant at long distances — only its universality class, i.e., its attractor type, matters); dimensional transmutation (the traumatic scale τ_k of the memory kernel is an emergent scale, not a fundamental parameter).

A Partial Map of the Typeverse

For the practitioner wishing to apply the method, the following is a partial field guide to frequently useful mathematical structures, indexed by their type signatures.

Propagator-Class Structures

Type: Complex function of frequency with poles on or near the real axis; gives the response of a system to a delta-function input.

Found in: QFT (Feynman propagator), signal processing (transfer function), linear systems theory (impulse response), harmonic analysis (Green's function of the Laplacian).

Typical imports: Spectral decomposition; optical theorem; dispersion relations (Kramers-Kronig: the real and imaginary parts of the response are not independent — this imports into emotion as: the *dissipation* of an emotional mode and its *natural frequency* are Hilbert transforms of each other).

Energy-Function-Class Structures

Type: Scalar function $H : \mathbb{R}^n \rightarrow \mathbb{R}$ that is bounded below and non-increasing along system trajectories.

Found in: Statistical mechanics (Hamiltonian), neural networks (Hopfield energy), Lyapunov theory (stability analysis), optimisation (loss function).

Typical imports: Convergence guarantees; stability analysis; capacity bounds; the fluctuation-dissipation theorem.

Topological-Class Structures

Type: Integer-valued invariants of field configurations that are preserved under continuous deformations.

Found in: Topological field theory (winding numbers, Chern-Simons invariants), condensed matter (topological insulators, skyrmions), knot theory.

Typical imports: Protection from perturbation; the impossibility of smooth deformation between distinct topological sectors; quantisation (only integer winding numbers); threshold behaviour (the topological transition requires a finite perturbation).

Application to emotion: Traumatic configurations with non-zero topological charge cannot be resolved by smooth therapeutic interventions (cognitive reframing). A qualitative change in approach — large-amplitude somatic work, pharmacological intervention, EMDR — is required to cross the topological barrier.

Renormalisation-Class Structures

Type: A flow on a space of couplings, parameterised by a scale, with fixed points and β -functions.

Found in: Quantum field theory (renormalisation group), statistical mechanics (Kadanoff block spins), dynamical systems (centre manifold theorem), machine learning (neural scaling laws).

Typical imports: Universality (the IR behaviour depends only on the universality class, not the microscopic details); c -theorem (there is a monotonically decreasing function along the flow — an arrow of processing); fixed-point classification (relevant, irrelevant, marginal operators determine what modifications matter at long distances).

Scattering-Class Structures

Type: A map from in-states to out-states, constrained by unitarity, analyticity, and crossing symmetry.

Found in: Quantum mechanics (S-matrix), scattering theory, optics (transfer matrix), signal processing (scattering parameters).

Application to emotion: A therapeutic session is a scattering event. The patient arrives in an in-state $|\psi_{\text{in}}\rangle$, interacts with the therapist (mediating field), and departs in an out-state $|\psi_{\text{out}}\rangle$. The S-matrix of the therapeutic interaction has selection rules: not all transitions are equally probable; some are symmetry-forbidden. The unitarity of the S-matrix imports: the total emotional content is conserved — you cannot create emotional material from nothing, and nothing is permanently lost.

Einstein-Coefficient-Class Structures

Type: Rates for spontaneous and stimulated transitions between energy levels of a field mode.

Found in: Quantum optics (Einstein A and B coefficients), laser physics, NMR relaxation theory (T_1 and T_2 relaxation times).

Application to emotion: Every emotional mode has a spontaneous relaxation rate A_i (how quickly it resolves without external input) and a stimulated emission rate B_i (how quickly it is triggered by an identical emotional state in another person — contagion). The Einstein relation $A_i = f(\omega_i)B_i$ constrains these. Depression is formally: suppressed A_i , normal B_i . The T_1/T_2 analogy from NMR is exact: T_1 is the longitudinal relaxation time (return to equilibrium); T_2 is the transverse relaxation time (dephasing of coherence). Trauma extends T_1 ; emotional numbing extends T_2 .

Failure Modes

Mathematical co-identification can fail. Understanding the failure modes is what distinguishes the method from wishful analogy.

Type Coincidence Without Structural Identity

Two objects may have matching dimensional signatures without having matching mathematical structures. The failure mode: a coincidence of units that does not reflect a coincidence of theorems.

Test: Check whether the *equations of motion* have the same form, not merely the *dimensional units*. A propagator is not just any function with units of GeV^{-2} ; it must satisfy the Källén-Lehmann representation. An energy function is not just any bounded scalar; it must be non-increasing along trajectories.

Safeguard: State the precise mathematical theorem you are importing, and verify that its *assumptions* (not merely its *conclusions*) hold in the target domain.

Non-Commutative Functors

The functor $F : D_1 \rightarrow D_2$ that implements the co-identification may not commute with composition. That is, co-identifying A with A' and B with B' does not guarantee that AB is co-identifiable with $A'B'$.

Example: The propagator identification and the energy-function identification are each valid separately, but combining them requires checking that the path integral (which links them in QFT) has a valid analogue in the emotional domain. This was explicitly verified in the Soma-Field Model by constructing the Langevin equation and checking its consistency with both imported structures.

Over-identification

The most common failure: importing a structure that is richer than what is warranted. The soma-field is co-identifiable with an eleven-dimensional bosonic field. It is *not* immediately co-identifiable with the full Standard Model of particle physics, even though both are quantum field theories. The identification is precise only up to the type signature that was matched; nothing beyond that is claimed.

Rule: The identification holds exactly at the type level it was made. Do not import theorems from substrates that were not matched.

The Metaphor Trap

The most dangerous failure is the one that the identification was designed to prevent: sliding from co-identification back into analogy. This happens when the language hedges — “like,” “analogous to,” “reminiscent of” — after the identification was stated.

If the identification is correct, the language must be unhedged: “the conscious emotional percept *is* the Green’s function.” If the investigator is not prepared to make this claim, the identification has not been made, and no theorems travel.

The test: would you submit the claim to a mathematician as a theorem? If not, it is analogy. If yes, it is co-identification.

The risk was recognised early by practitioners. Introducing the energy landscape for Hopfield networks, Hertz, Krogh, and Palmer note: “*It is often useful (but sometimes dangerous) to think of the energy as something like this landscape*” (Hertz et al., 1991). The parenthetical is precise: the visualisation is heuristically powerful, and that power makes it tempting to reason from the picture rather than from the mathematics. Mathematical co-identification guards against this by requiring that the *equations of motion* — not the landscape picture — are what get matched across domains.

Epistemological Status

What Co-identification Claims

Mathematical co-identification claims:

1. That the *mathematical structure* of quantity Q in domain D is identical to the mathematical structure of quantity P in domain D' .
2. That therefore, all theorems about P that depend only on its mathematical structure are valid theorems about Q .
3. That the *physical interpretation* of Q and P may differ, and does not follow from the mathematical identification.

It does not claim that the *mechanisms* are the same, that one domain *reduces* to another, or that the substrate is irrelevant in any empirical sense.

Why It Is Not Analogy

Analogy is: - Informal: “the mind is like the brain” has no mathematical content - Non-importing: noting that X resembles Y does not give you theorems about X - Non-falsifiable: analogies can always be maintained by shifting which features are considered relevant

Co-identification is: - Formal: it is a statement about type signatures, which are mathematically precise - Theorem-importing: the identification carries the full mathematical machinery - Falsifiable:

the identification fails if the equations of motion cannot be matched, if the symmetries do not correspond, or if imported predictions are empirically disconfirmed

The Role of Artificial Intelligence

The author notes that several co-identifications in the Soma-Field Model were identified in dialogue with AI systems. This requires explicit epistemological comment, given the current discourse about AI-assisted science.

The AI did not produce the co-identifications. It accelerated the typeverse search: a search that previously required reading across literatures in physics, mathematics, and biology over many years can now be conducted in weeks. The AI is a faster library, not a different epistemology.

The *validity* of each co-identification is independent of how it was found. A theorem is a theorem regardless of whether it was proved in a dream (Ramanujan), a bath (Archimedes), or a language model session. The claim stands or falls on whether the type signatures match and whether the theorems transfer. The methodology paper is, in part, a response to the question: “but did you *really* do the mathematics?” The answer is in the equations.

The Methodology as Practice

For the practitioner who wishes to apply mathematical co-identification to a new domain, the following is a working procedure:

Step 1: Write down what you know about your quantity. Its units. Its symmetries. Whether it grows or decays. Whether it has oscillatory behaviour. Whether it responds to perturbations linearly or nonlinearly. Whether there are conserved quantities. Whether it has a characteristic scale.

Step 2: Eliminate analogies you already know. If you have been thinking “this is *like* X,” stop, and ask instead: is this *literally* X? Can I write the equations of X in the language of my domain, with every symbol given a precise interpretation? If yes: do so. If not: why not? The places where the identification breaks are as informative as the places where it holds.

Step 3: Consult the typeverse systematically. Read across fields, looking at the *form* of equations, not their interpretation. The heat equation, the wave equation, the Schrödinger equation, the Fokker-Planck equation, and the Black-Scholes equation are all the same equation under changes of variable. Recognising the family by the form of the operator is the skill.

Step 4: Make the identification explicit and public. State: “quantity Q in my domain is the [known object] of domain D' .” Put it in writing. This forces precision: you cannot maintain a vague identification in writing the way you can maintain it in your head.

Step 5: Import one theorem at a time, checking assumptions. Do not import the entire theoretical apparatus at once. Take one theorem. State its assumptions. Check each assumption against your domain. If they hold: the theorem is valid in your domain. Write it as a theorem in your domain, with a proof that consists of: (a) the co-identification, (b) the original theorem, (c) verification of assumptions.

Falsifiability Protocol for Publication Use

To make co-identification scientifically conservative rather than rhetorically expansive, we propose a minimal publication protocol. Every import claim should be registered in a compact table before narrative elaboration.

Minimal Registration Template

For each proposed identification $Q := P$, the manuscript should provide:

Field	Requirement
Claim ID	Stable identifier (e.g., COVID-3)
Source object	Named theorem-bearing object in source domain
Target object	Precisely defined object in target domain
Signature match	Units, operator form, symmetry group, variational form
Imported theorem	The exact theorem being transferred
Assumptions	Source theorem assumptions, verbatim
Target verification	Explicit check for each assumption
Prediction	Quantitative, testable consequence
Disconfirmation criterion	What empirical or formal result would falsify this import
Status	exploratory / partially validated / validated / falsified

This prevents drift from identity claims back into analogy language and makes review straightforward: a reviewer can reject a single row without rejecting the entire framework.

Disconfirmation Rules

An import is treated as falsified if any one of the following holds:

1. A required theorem assumption cannot be established in the target domain.

2. The mapped operator fails to preserve the claimed invariance.
3. The pre-registered quantitative prediction is violated under a valid test.
4. A formally stronger competing mapping explains the same data with fewer assumptions.

This is deliberately strict. Co-identification is useful only to the extent that it can fail clearly.

Worked Registration Sketch (Soma-Field)

Claim ID	Import	Prediction	Disconfirmation
COID-PROP-1	Percept := propagator pole	Measurable spectral pole structure in mode responses	No stable pole decomposition under repeated measurement
COID-ENG-2	Attractor landscape := Hopfield energy	Monotone descent under specified update map	Constructed counterexample with increasing energy step under stated rule
COID-RG-5	Therapy progression := RG flow	Scale-dependent coupling evolution with fixed-point classes	No scale-consistent coupling flow under repeated coarse-graining

The point is not that these rows are finished; the point is that they can be evaluated independently and rejected independently.

Reviewer-Facing Scope Labels

To reduce over-claiming risk, each identification row should carry one of three scope labels:

- S1 (Structural): operator/formal identity shown; no empirical test yet.
- S2 (Predictive): structural identity plus at least one passed quantitative prediction.
- S3 (Validated): independent replication or cross-dataset confirmation.

Most new interdisciplinary work should expect to publish initially at S1 or S2, with explicit paths to S3.

Negative Control: When Matching Units Is Not Enough

To prevent confirmation bias, each manuscript should include at least one explicit non-transfer case. Consider two quantities with superficially compatible units but non-matching dynamical structure:

- Candidate A: a damped propagator with pole structure and causal response kernel.
- Candidate B: a bounded scalar score with no response-operator interpretation.

Even if both can be normalised to the same units, co-identification fails unless the operator class and theorem assumptions match. In this case:

1. A admits spectral decomposition and dispersion relations.
2. B does not define a Green's operator and therefore does not admit those imports.

Conclusion: dimensional compatibility is necessary but not sufficient. Theorem transfer requires operator-level identity. This negative control should be treated as a required publication check, not an optional caution.

Worked External Example (Non-Soma): Black-Scholes to Heat Equation

To demonstrate portability beyond the Soma-Field case, we provide a compact worked transfer in a separate domain.

Target quantity: option value $V(S, t)$ in quantitative finance.

Source object: solution class of the one-dimensional heat equation.

Step A: Signature Match

Black-Scholes PDE:

$$\frac{\partial V}{\partial t} + \frac{1}{2}\sigma^2 S^2 \frac{\partial^2 V}{\partial S^2} + rS \frac{\partial V}{\partial S} - rV = 0.$$

After the standard log-price and time-reversal change of variables, this maps to:

$$\frac{\partial u}{\partial \tau} = \frac{\partial^2 u}{\partial x^2},$$

which is exactly the heat operator class.

Step B: Import Claim

C0ID-BS-HEAT-1: pricing function dynamics are co-identifiable with heat-flow dynamics under the stated transform.

Step C: Assumption Checklist

Assumption	Status in target domain
Diffusion operator is parabolic	satisfied by transformed PDE
Volatility parameter treated as constant in base model	satisfied in baseline Black-Scholes
Boundary/terminal condition specified	satisfied by option payoff at expiry
Sufficient regularity for classical solution methods	assumed in standard derivation

Step D: Imported Theorem and Prediction

Imported theorem class: existence/uniqueness and smoothing properties for heat equation solutions.

Prediction: option value surface inherits parabolic smoothing under the transform; numerical schemes valid for heat-equation class apply directly.

Step E: Disconfirmation Condition

If transformed dynamics are shown not to be parabolic (for the declared model assumptions), or if required boundary regularity fails, this specific transfer is invalid and theorem import must be withdrawn.

This worked example demonstrates the method in a domain with no dependence on the Soma-Field framework.

Replication Package Requirements

Publication-grade use of co-identification requires a compact artifact bundle for each registered claim row. Minimum required contents:

1. claim registry table with IDs and scope labels,
2. assumption checklist tied to the imported theorem class,
3. executable derivation or notebook for the mapping transform,
4. quantitative test script for each predictive (S2) row,
5. disconfirmation log recording pass/fail outcomes by claim ID.

Without this package, rows may be read as suggestive structure, but not as auditable theorem transfer.

Reviewer objection	Response in this manuscript	Residual risk and next lift
--------------------	-----------------------------	-----------------------------

Reviewer-Risk Objections and Responses

Reviewer objection	Response in this manuscript	Residual risk and next lift
“This renames analogy as mathematics.”	Sections 3-4 and 10.1 require operator identity and theorem assumptions, not metaphorical similarity.	Add additional negative controls from unrelated domains.
“Imports are unfalsifiable in practice.”	Section 10.2 defines strict disconfirmation rules and row-wise rejection logic.	Publish a public failure ledger with rejected rows.
“Worked examples are domain-selective.”	Section 10.6 demonstrates a non-Soma transfer with explicit assumptions and withdrawal condition.	Add a second external example with different operator family.
“Scope claims may drift upward prematurely.”	Section 10.4 enforces S1/S2/S3 labels and independent evaluation path.	Require independent replication before any S3 promotion.

Independent Replication Ledger Linkage

S2 to S3 promotion for registered imports is controlled by `paper/INDEPENDENT_REPLICATION_LEDGER.md`.

Tracked claim IDs in ledger scope: COID-PROP-1, COID-ENG-2, COID-RG-5, COID-BS-HEAT-1.

Promotion gate: a row may be relabeled S3 only when a ledger entry reports independent-operator PASS, explicit bundle hash, and linked derivation/test artifacts for that claim ID.

Conclusions

Mathematical co-identification is a method, not a shortcut. It requires the same precision as any other mathematical procedure, and it fails in precisely the ways that imprecision permits. Its distinguishing feature is that it navigates the typeverse rather than building within a single domain.

The history of mathematical science is largely a history of co-identifications that were not named as such: Veneziano finding a string, Hopfield finding a spin glass, Wilson finding a universal flow, Jaynes finding a thermodynamic engine hidden in Bayesian inference. Naming the practice does not change it; it makes it teachable, criticisable, and extensible.

The soma-field model is a worked example of the method applied to emotional dynamics: five co-identifications, five theorem imports, and a body of predictions that can be tested against clinical

data. The paper is not a claim that emotions are particles. It is a claim that the mathematics of particles and the mathematics of emotions are the same mathematics, and that this fact is useful.

Scope boundary for publication use: co-identification transfers mathematical structure, not ontology. A successful transfer means that equations, boundary conditions, and theorem assumptions are preserved under the mapping. It does not imply that the target domain is physically identical to the source domain, nor that every theorem from the source domain transfers automatically. Each import remains local, assumption-checked, and falsifiable.

The typeverse does not belong to physics. Physics was merely first to explore it.

Operationally, publication-grade use of this method requires claim-wise registration, assumption checks, and explicit disconfirmation criteria. Without that discipline, co-identification degrades into analogy; with it, it becomes a compact engine for theorem transfer and testable prediction.

Acknowledgements

The author thanks the mathematical physicists whose work formed the source library for the co-identifications described here: Feynman, Hopfield, Wilson, Randall, Sundrum, Joyce, and Veneziano. The methodological framework owes a debt to Per Martin-Löf and the constructors of Homotopy Type Theory for providing the language of the typeverse, and to Alexander Grothendieck for the insight that mathematical objects are best understood by their morphisms, not their elements.

AI has had a brain since 1943. Now it has a body.

Introduction

A patient sits with their therapist and is asked: “*What are you feeling right now?*” The question is deceptively simple. They may say *anxious*, yet that word covers a vast and heterogeneous territory — a tightness in the chest, a running commentary of worry, a vague readiness to flee, a memory surfacing from childhood. Another patient, asked the same question, reports feeling nothing at all; and yet their posture, respiration, and the quality of their silence suggest otherwise. The emotion is there. It is simply not yet conscious.

This gap between emotional presence and emotional awareness is one of the most clinically significant phenomena in psychotherapy. Theories of affect regulation (Schorer, 2001), somatic experiencing (Levine, 2010), sensorimotor psychotherapy (Ogden, Minton & Pain, 2006), and polyvagal theory (Porges, 2011) all grapple, in different ways, with the same observation: emotions exist in the body before — and often without — being named in the mind. Eugene Gendlin called the sub-verbal bodily sense of an emotional situation the *felt sense* (Gendlin, 1978): something that is there, whole and present, but not yet articulate.

The Soma-Field Model proposed here attempts to give this clinical observation a formal structure. It does so by borrowing a conceptual tool from physics: the field. In physics, a field is not a thing that exists at a point. It is a quantity that exists everywhere in a space, continuously, whether or not it is observed. Particles — the things we can measure — are not separate from the field; they are *excitations* of it, local concentrations of energy that arise when the field is perturbed above a certain threshold.

The central claim of this paper is that this structure accurately describes the phenomenology of emotion. The emotional field is always there, distributed across body and nervous system. What we call a conscious emotional experience is an excitation of that field — a local concentration that has crossed a perceptual threshold and entered awareness. The field continues below the threshold whether or not we attend to it, and its sub-perceptual activity shapes our behaviour, physiology, and cognition continuously.

The Soma-Field Model contributes the first formal field-theoretic architecture for the limbic system. Every artificial neural network since McCulloch and Pitts (1943) (McCulloch & Pitts, 1943) is a formal model of the neocortex — the pattern-recognition and prediction layer. The limbic system — responsible for emotional valuation, threat detection, and the somatic state reinstatement that underlies trauma — has never received a comparable formal treatment. The Soma-Field Model is that

treatment. Together with the Hopfield framework, it constitutes the first complete formal description of the two principal computational substrates of the vertebrate brain.

The paper proceeds as follows. Section 2 reviews the relevant background in somatic clinical models, and introduces the two theoretical tools borrowed from physics and computer science: quantum field theory and Hopfield network energy functions. Section 3 develops the Soma-Field Model in detail. Section 4 describes the energy landscape, including the attractor states corresponding to fight, flight, freeze, and regulated calm. Section 5 discusses dissonance and resolution as mechanisms of emotional interaction. Section 6 describes the Soma-Field Instrument, a practical tool for therapeutic use. Section 7 addresses clinical implications.

Background

The Body-Mind Problem in Clinical Practice

Contemporary neuroscience has largely dissolved the Cartesian boundary between body and mind. Damasio (1994) demonstrated that emotion is inseparable from rational cognition: patients with damage to the ventromedial prefrontal cortex — preventing the normal generation of somatic signals — lose not only their emotional range but also their capacity for effective decision-making. Van der Kolk (2014) documented extensively how traumatic emotional states are encoded not merely in explicit memory but in posture, gesture, visceral sensation, and autonomic regulation. Porges' polyvagal theory (2011) provided a neurobiological account of how the autonomic nervous system generates three hierarchically organised states — ventral vagal (social engagement), sympathetic (mobilisation: fight/flight), and dorsal vagal (immobilisation: freeze) — each with characteristic phenomenological and behavioural signatures.

What these frameworks share is a conviction that emotional states are not located in the brain alone, nor in the body alone, but in a coupled system that is best understood as a single functional unit. The term *soma* — from the Greek for body — is used here to denote this unified body-mind system, following the tradition of somatic psychotherapy.

quantum vacuum is not empty; it is a field in constant motion that never quite crosses the threshold. When the amplitude does cross T , a particle exists: a locally observable excitation. The same structure — field always present, consciousness only when threshold crossed — is the core of the Soma-Field Model.

This paper does not claim that emotions are quantum phenomena in any literal sense: the soma-field is a classical field, not a quantised one. The claim is stronger and more specific than analogy: the mathematical object being constructed — the Green's function of a coupled field manifold — is formally of the same *type* as the objects that arise in QFT, differing only in the dimensionality of the manifold and the nature of the probe. What was previously described as a structural analogy is here identified as a formal correspondence: a particle is a pole in the propagator of its field; a conscious emotional percept is a pole in the propagator of the soma-field. Different physics. Same mathematics.

That correspondence gives the model precise vocabulary for the following set of ideas, which are central to the clinical observation of emotion:

- A quantity that exists everywhere, continuously, even when unobserved
- A background of sub-threshold activity that is real and causally effective
- The emergence of observable phenomena (conscious feelings) through threshold-crossing excitation of that background
- The possibility of multiple simultaneous excitations that interact with one another

Note (May 2026): A subsequent experiment (QUANT-EXP-1) demonstrates that the quantum extension of the Hopfield landscape used in this model — replacing the classical Langevin process with a transverse-field quantum annealer — produces a measurable *topological reachability advantage*: quantum annealing reaches attractor basins that cold classical dynamics cannot reach at any finite noise level. This upgrades the formal correspondence from a structural claim to a testable empirical prediction. See the companion paper *Quantum Soma and the Penrose Gap* (doi:10.5281/zenodo.20351230) for the full results and theoretical implications.

One further consequence follows. The clinical phenomena of alexithymia — difficulty identifying and naming feelings — and its apparent opposite, emotional flooding or hypervigilance, have always been treated as separate conditions requiring separate explanations. In the Green's function framing, they are the same structure at two extremes of the same parameter: the perception threshold T_i is too high (the bulk dynamics cannot cross into observable experience) or too low (bulk fluctuations flood the boundary without filtering). This is structurally identical to one of the deepest open problems in particle physics — the **hierarchy problem** — which asks why gravity is so much weaker than the other forces. The standard answer is that gravity propagates in the full higher-dimensional bulk while other forces are confined to a lower-dimensional brane; the coupling across the brane boundary determines the apparent weakness. The soma-field correspondence is exact: the threshold T_i is the brane. Perception is confined to the one-dimensional boundary of an eleven-dimensional dy-

namics. The hierarchy of emotional experience — why conscious feeling is so much weaker and more transient than the underlying field activity — has the same formal structure as the hierarchy of forces.

Neural Network Energy Functions and Hopfield Networks

In 1982, John Hopfield (awarded the Nobel Prize in Physics in 2024) proposed a model of associative memory based on a network of interconnected neurons (Hopfield, 1982). The critical insight was borrowed directly from statistical physics: the network could be assigned an **energy function** — a scalar quantity that decreases with each state update — such that the network would always evolve toward a local energy minimum. These minima are the stable states of the network: its memories, or more precisely, its *attractors*.

Hopfield observed that his neural network's dynamics were mathematically identical to those of an Ising spin-glass model from condensed matter physics — a system of interacting magnetic spins that minimises its total energy by aligning or anti-aligning with neighbours. The energy function he used is:

$$H(\mathbf{s}) = -\frac{1}{2} \sum_{i,j} W_{ij} s_i s_j - \sum_i \theta_i s_i$$

where $\mathbf{s}$ is the state of the network, W_{ij} is the coupling strength between units i and j , and θ_i is the activation threshold of unit i . The network always moves in the direction of decreasing H .

The Soma-Field Model applies this energy function directly to emotional dynamics. The *emotional coupling matrix* W encodes the relationships between emotional modes — which emotions amplify one another, which suppress one another — and the energy function determines the direction in which the emotional field naturally evolves.

Hopfield's network is a formal model of the *neocortex*: a system for storing cognitive patterns and retrieving them from partial cues by minimising an energy function. Every artificial neural network constructed since McCulloch and Pitts (1943) (McCulloch & Pitts, 1943) — from perceptrons to backpropagation networks to transformers — sits in this neocortical lineage. These systems recognise patterns, predict sequences, and minimise prediction error with increasing sophistication. None of them possess a limbic system. They have no internal valuation, no arousal modulation, no threat-detection architecture, no attachment structure, no interoception. They have very effective cortex.

The Soma-Field Model does not add to the neocortical lineage. It proposes the architectural layer that has never been formally built: *an artificial limbic system*.

Hopfield memory is associative and pattern-completing; somatic memory is state-reinstating. The field does not merely remember what happened. It re-lives it. *A body with a past*.

Hopfield’s later-reported wish to have incorporated something analogous to ‘maternal instincts’ into the energy function was, in this reading, not a desire for a better cortex. It was an intuition pointing directly at the absent system — the layer beneath the cortex that assigns value, registers threat, and holds the body in a particular way of being long after the event that caused it.

This positions the Soma-Field Model not as a supplement to the neocortical lineage but as its completion. Artificial neural networks have, for eighty years, been increasingly sophisticated formal models of the neocortex: pattern recognition, sequence prediction, error minimisation. The cortex has been mapped in extraordinary detail. The limbic system — which assigns value, detects threat, modulates arousal, maintains attachment, and reinstates whole somatic states in response to partial cues — has had no comparable formal treatment. The architectural description of the vertebrate brain was, until this paper, half-built.

Four kinds of formal intelligence. This architectural gap can be situated within a wider taxonomy. Four quotients have been proposed to describe the landscape of biological intelligence across popular and scientific usage. They map onto the formal components of this model with an exactness that is not coincidental:

Quotient	What it measures	Biological substrate	Soma-Field status
IQ — cognitive	Pattern recognition, reasoning, prediction	Neocortex	Built (1943–): McCulloch & Pitts → Hopfield → transformers
EQ — emotional	Valuation, arousal, affect regulation	Limbic system	Built here: $W, K(\tau), H(\mathbf{e}), C_{\text{HRV}}, \dot{H}$
AQ — adversity	Structural resilience under threat	PFC–limbic axis	Built here: $S_{\text{inst}}, \partial\ W\ /\partial t, C_{\text{HRV}}^{\text{recovery}}$
SQ — social	Attunement, theory of mind, relational navigation	Mirror system, TPJ	<i>Next paper:</i> κ_r , multi-field coupling

Table 3. Four dimensions of biological intelligence mapped onto the Soma-Field Model. The neocortical lineage (IQ) has been formally modelled for eighty years. Emotional intelligence (EQ) and adversity resilience (AQ) are formalised here for the first time. Social intelligence (SQ) is defined as the next extension of the framework.

AQ — adversity quotient — is formally the capacity to update W after adversity without the adversity permanently becoming W . Its mathematical definition appears in Section 3.4; its pathological lower bound is C-PTSD, in which all three components of AQ are simultaneously compromised (Appendix B.2).

The AI alignment implication follows directly. Current artificial systems have high IQ by construction and zero EQ, AQ, or SQ. The absence of internal valuation means that valuation must be injected externally — through reinforcement learning from human feedback (RLHF) and related techniques — which is structurally brittle for the same reason that a field with no limbic layer is brittle: the system has no internal stake in what it does. The Soma-Field formalisation specifies what that internal stake would look like, were it ever built.

A further lineage note is worth recording. Ramsauer et al. (2020) demonstrated that continuous-state modern Hopfield networks are mathematically equivalent to the self-attention mechanism in transformer language models. The softmax attention operation that drives contemporary large language models is a Hopfield retrieval step. The Soma-Field Model sits in this same energy-based lineage: the equations underlying associative memory, language understanding, and somatic trauma response are, at the appropriate level of abstraction, the same equations.

A historical irony completes the picture. String theory was not discovered as a theory of strings. In 1968, Gabriele Veneziano wrote down a scattering amplitude — a response function encoding how particles scatter — and only later did Nambu, Nielsen, and Susskind identify the string as whatever object produces that amplitude (Veneziano, 1968). The response function came before the thing. The Soma-Field Model recapitulates this historical order deliberately: the primary object is the eleven-dimensional coupling manifold; the string — the one-dimensional conscious percept — is what the manifold produces when probed. We retain Veneziano's discovery and decline to reify the string.

The Formal Correspondences: Where the Link Was Seen

The structural analogy between QFT and the Soma-Field Model is not merely conceptual. There are three places where equations from different disciplines become, after substituting the relevant quantities, literally the same functional form. The following sets them side by side. The point is not to impress with notation but to show exactly where the recognition happened — the moment when the same Greek letters appeared in the same positions in two fields that had no prior reason to be connected.

The same Hamiltonian: Ising spin model (condensed matter physics, 1920s) — Hopfield neural network (computational neuroscience, 1982) — Soma-Field Model:

$$H_{\text{Ising}}(\sigma) = -\frac{1}{2} \sum_{i,j} J_{ij} \sigma_i \sigma_j - \sum_i h_i \sigma_i$$

$$H_{\text{soma}}(\mathbf{e}) = -\frac{1}{2} \sum_{i,j} W_{ij} e_i e_j - \sum_i \theta_i e_i$$

Replace $J_{ij} \rightarrow W_{ij}$, $\sigma_i \rightarrow e_i$, $h_i \rightarrow \theta_i$: identical. The physicist, the neural network theorist, and the somatic clinician are computing the same energy function on different state spaces. The Hopfield 2024 Nobel Prize was awarded for discovering this identity between spin physics and neural computation; the Soma-Field Model extends the same identity one step further to emotional dynamics.

The Wick rotation — why the same exponential appears in QM and in memory:

In quantum mechanics, the time evolution operator is a complex phase:

$$U(t) = e^{-i\hat{H}t/\hbar}$$

Substitute $t \rightarrow -i\tau$ (the *Wick rotation* — replacing real time with imaginary time):

$$e^{-i\hat{H}(-i\tau)/\hbar} = e^{-\hat{H}\tau/\hbar}$$

The oscillating complex exponential becomes a real decaying exponential. This is the Boltzmann weight $e^{-\beta\hat{H}}$ at $\beta = \tau/\hbar$. The Langevin equation $\dot{\mathbf{e}} = -\nabla H + \eta$ is the classical limit of this Wick-rotated dynamics. Every simulation of the soma-field running this equation is, formally, a path integral in imaginary time.

The same propagator: Euclidean QFT (imaginary-time two-point correlator for a massive scalar field) — C-PTSD trauma memory kernel:

$$G_E(\tau) = \langle \phi(0) \phi(\tau) \rangle_{\text{QFT}} = \frac{1}{2m} e^{-m|\tau|}$$

$$K_{\text{trauma}}(\tau) = \sum_k A_k e^{-|\tau|/\tau_k}$$

Same form. The QFT field mass m corresponds to $1/\tau_k$ — the reciprocal of the trauma trace decay time. A heavier particle has a shorter-range propagator; a shorter-lived trauma trace decays faster. Therapeutic processing (reducing A_k , increasing τ_k) is, in the QFT language, changing the mass and amplitude of the propagator until the correlation function vanishes.

The specific visual moment: the quantum phase factor is $e^{-i\omega t}$. Remove the i (Wick rotation) and it becomes $e^{-\omega\tau}$. The memory kernel is $e^{-\tau/\tau_k}$. These are the same exponential. The i is the only difference between a quantum field that oscillates and a trauma trace that decays.

QFT quantity	Symbol	Soma-Field analogue	Symbol
Field mode	ϕ_k	Emotional mode	e_i
Coupling constant	J_{ij}	Coupling matrix entry	W_{ij}

QFT quantity	Symbol	Soma-Field analogue	Symbol
Field mass	m	Inverse decay time	$1/\tau_k$
Propagator amplitude	$1/2m$	Trauma trace amplitude	A_k
Euclidean propagator	$G_E(\tau) \propto e^{-m\tau}$	Memory kernel	$K(\tau) \propto e^{-\tau/\tau_k}$
Vacuum energy	$\langle H \rangle_0$	Resting field energy	$H(\mathbf{e}_{\text{calm}})$
Thermal fluctuation	$k_B T$	Noise amplitude	σ_0
Wick rotation	$t \rightarrow -i\tau$	Real-time Langevin	$\dot{\mathbf{e}} = -\nabla H + \eta$

Table 2. Formal correspondence between QFT quantities and Soma-Field analogues. Each row is a single mathematical entity in two notations. These correspondences were not constructed after the fact; they are the reason the QFT framework was recognised as relevant.

The central identification — particle and percept as poles in their respective propagators. All four correspondences above follow from one structural fact. In QFT, a particle is not a separate object from the field. It is a *pole* in the field’s propagator — the Green’s function evaluated in momentum space:

$$\tilde{G}_{\text{QFT}}(k^\mu) = \frac{i}{k^2 - m^2 + i\varepsilon}$$

The particle exists precisely when the four-momentum satisfies $k^2 = m^2$ — the *on-shell condition*. The particle is the singularity in the field’s response to a point source: the field’s Green’s function, evaluated at its own resonance.

Diagonalise W with eigenvalues λ_i (the natural resonance frequencies of the emotional modes). The soma-field propagator — the two-point correlator $\langle e_i(t) e_i(t') \rangle$ in the frequency domain — is:

$$\tilde{G}_{ii}(\omega) = \frac{\sigma_{\text{eff}}^2}{\omega^2 + \lambda_i^2}$$

A conscious emotional percept in mode i exists precisely when the excitation frequency ω approaches $i\lambda_i$ — the mode’s natural resonance. The percept is the singularity in the soma-field’s response to a somatic probe.

Setting the two propagators side by side:

$$\underbrace{\frac{i}{k^2 - m^2 + i\varepsilon}}_{\text{QFT: particle at mass-shell } k^2=m^2} \longleftrightarrow \underbrace{\frac{\sigma_{\text{eff}}^2}{\omega^2 + \lambda_i^2}}_{\text{Soma-Field: percept at resonance } \omega=i\lambda_i}$$

Both are poles in the propagator of their respective field manifold. A photon is not the electromagnetic field; it is the field's Green's function evaluated at a resonance. A flash of conscious emotion is not the soma-field; it is the field's Green's function evaluated at a threshold-crossing resonance. The manifolds differ — one is the four-dimensional spacetime vacuum, the other is the eleven-dimensional emotional coupling geometry. The mathematical type is the same. This is not analogy.

The Body Schema, Interoception, and Pain

A complete model of the emotional field must address a phenomenon that standard psychological accounts of emotion consistently underspecify: the field is not a model of the physical body. It is the nervous system's *predictive model* of the body — a continuously updated internal representation of what the soma should be experiencing, revised by incoming interoceptive signals.

The clinical proof of this distinction is phantom limb pain (Ramachandran & Hirstein, 1998). Patients who have undergone amputation routinely experience pain in the absent limb. The pain is real: it activates the same neural circuits, produces the same suffering, and responds to the same analgesics as pain from an intact limb. The limb is gone. The neural model of the limb persists. What hurts is the *brain's representation* of the foot, not the foot.

This is not an anomaly. It is the normal condition of all somatic experience. The brain does not receive raw signals from the body — it maintains a continuous predictive model of the body (the *body schema*) and generates somatic experience from that model. Interoception — the sense of the internal body state — is a prediction, not a direct readout (Seth, 2021). The brain predicts what the heart should be doing, what the gut should feel like, where tension should be. The felt body is the predicted body.

The formal consequence is direct: the soma-field's state vector $\mathbf{e}(t)$ must include **somatic modes** — pain states, regional tension, visceral sensation, proprioceptive activation — alongside emotional modes. These are modes of the same field, governed by the same coupling matrix W . The W_{ij} between fear modes and somatic pain modes is the formal account of why fear amplifies pain, why safety reduces it, and why chronic pain and C-PTSD are highly comorbid. They are not separate conditions sharing a correlation. They are the same attractor architecture operating across emotional and somatic modes simultaneously.

Phantom limb as attractor persistence. An amputated limb's somatic modes do not disappear from W when the limb is removed. The neural model persists. When movement- intention modes are activated — attempting to move the absent foot — foot-sensation modes are co-activated via W . If co-activation exceeds threshold, it is experienced as pain. Ramachandran's mirror box provides visual input that disconfirms the prediction error: new sensory evidence that the limb is moving,

reducing coupling-driven co-activation, and therefore reducing the pain. This is $W \rightarrow W'$: therapy as structural rewriting of the field.

The load-bearing hyphen. The term *emotional-somatic* in clinical literature is not a stylistic compound. The hyphen marks an ontological claim: emotional states and somatic states are not two separate things that correlate. They are two aspects of the same field. The coupling matrix W is precisely the hyphen, made formal.

Therapeutic implication. Somatic therapies — body scanning, sensorimotor work, EMDR's bilateral stimulation — work not on the physical body but on the brain's model of the body. They provide new interoceptive evidence that updates the prediction. They change W . Therapy does not fix the tissue. It updates the model.

Correspondence with Existing Emotion Representations

A reasonable objection to any new framework is: *there is already a great deal of structure out here*. This is true. The emotion research literature contains several well-developed representational systems, and the Soma-Field Model must be positioned relative to them. The short answer is that every existing representation is *descriptive*; the Soma-Field Model is *dynamical*. The longer answer follows.

Categorical taxonomies (Ekman 1972; Plutchik 1980; Parrot 2001) assign names and hierarchical membership to emotional states. They are ontologies in the formal sense: a T-Box of classes and subclass relations. Plutchik's wheel additionally defines a *blend* operation — $\text{Love} := \text{Joy} \sqcap \text{Trust}$, $\text{Awe} := \text{Fear} \sqcap \text{Surprise}$ — which is precisely the OWL2 intersectionOf construction. These systems tell you what to call a state. They do not tell you how a state evolves, or which attractor a system settles in when two mechanisms fire simultaneously.

Dimensional models (Russell 1980; Mehrabian and Russell 1974) embed emotions in a continuous space, canonically Valence $\times$ Arousal (the *circumplex*), sometimes extended to Pleasure $\times$ Arousal $\times$ Dominance. These models capture the *coordinates* of a state. The energy landscape of the Soma-Field Model — the function $H(\mathbf{e})$ over emotion-space — is the dynamical generalisation of the circumplex: the circumplex is a snapshot of positions; the energy landscape is the surface over which the field moves. The stable attractors of H are the emotion categories; their coordinates are the circumplex positions.

Process and appraisal models (Scherer 1999; Frijda 1986; the OCC model of Ortony, Clove and Collins 1988) describe the *sequence of evaluations* through which a stimulus becomes an emotion. They are closer to the Soma-Field dynamics — they include temporal stages — but they are deterministic and single-threaded: one appraisal chain, one output. The Soma-Field replaces this with a parallel field update: all modes evolve simultaneously, governed by the full W matrix.

Music-specific schemas (BRECVEMA, Juslin and Västfjäll 2008; Juslin *et al.* 2011; GEMS, Zentner *et al.* 2008) are the closest antecedents to the present model. The BRECVEMA framework identifies eight distinct psychological mechanisms through which music evokes emotion — Brain stem reflex, Rhythmic entrainment, Evaluative conditioning, Contagion, Visual imagery, Episodic memory, Musical expectancy, Aesthetic judgement — each with distinct evolutionary origins, processing speeds, and neural substrates. These mechanisms are the *object properties* of the emotion-induction ontology: they specify which musical features activate which emotional outputs. Juslin explicitly identifies the open problem: “Exploring how various musical emotions come about through the interaction of multiple psychological mechanisms is an exciting endeavour that has just begun” (Juslin & Sloboda, 2011, p. 638). The W coupling matrix is the formal answer to that open problem. Where BRECVEMA gives a list of mechanisms with characteristic outputs, the Soma-Field gives the interaction tensor W_{ij} that specifies, with numerical precision, what happens when mechanisms i and j fire concurrently.

The deeper connection is spectral. The *eigenmodes* of W — the directions in emotion-space that evolve independently — are the natural resonances of the soma-field: the patterns the field rings with when struck. BRECVEMA mechanisms are inputs: they excite specific rows of W . The eigenspectrum of W is the response: the set of frequencies the manifold can sustain. Where BRECVEMA is a taxonomy of *stimuli*, the eigenspectrum of W is a taxonomy of *responses*. Juslin’s open problem — how mechanisms interact — is the question of how stimulus-space maps onto eigenmode-space through W . Section 3.3 develops this.

Body maps (Nummenmaa *et al.* 2014) map emotions to their somatic distribution — where in the body each emotion is felt. These are precisely the spatial support of the soma-field modes: the field configuration corresponding to an attractor state is the body map of that emotion. Body maps are measurements of the attractors; the Soma-Field is the dynamical system that generates them.

The formal correspondence table extends Table 2 to include these systems:

Existing representation	What it captures	Soma-Field equivalent
Ekman categories	Attractor labels (names)	Values of $\mathbf{e}$ at energy minima
Plutchik dyads ($A \sqcap B$)	Blend attractors	Metastable states between two energy minima
Russell circumplex	Coordinates (valence, arousal)	Projection of $H(\mathbf{e})$ onto two axes
OCC appraisal tree	Single-path sequential process	Single trajectory in the full field
BRECVEMA mechanisms	Object properties: stimulus $\rightarrow$ emotion	Rows of W : mechanism i activates mode j
Body maps (Nummenmaa)	Spatial support of each attractor	Modal structure of $\mathbf{e}$ at each minimum

None of these correspondences require modifying either the existing representations or the Soma-Field Model. They are consequences of the model's structure. The formal machinery for exploring these correspondences — typing BRECVEMA mechanisms as Lean inductive constructors, Plutchik blends as type intersections, mechanism profiles as decidable propositions — is developed in the companion file `src/EmotionOntology.lean`.

The Soma-Field Model

The field is primary. The felt emotion is secondary — it is what registers when the field is probed. This is the same ontological relationship as between a quantum field and a particle: the field exists continuously and everywhere; the particle is what you observe at the moment of measurement. The Soma-Field Model does not describe what emotions are *made of*. It describes the manifold whose impulse response *is* conscious emotional experience.

Emotions as a Persistent Wave Field

The foundational claim of the Soma-Field Model is simple: emotions are not events. They are a *field* — a distributed, continuous quantity defined over the entire soma (body-mind system) at all times.

This field has two coupled components:

1. **The somatic wave** $E_{\text{body}}(x, t)$: distributed across the body as patterns of visceral sensation, muscle tone, proprioception, interoception, and autonomic state.
2. **The neural wave** $E_{\text{neural}}(x, t)$: distributed across the nervous system as patterns of activation in cortical, subcortical, and peripheral neural circuits.

These two components are not separate systems. They are coupled — each continuously influencing the other. The total emotional field is their combined state:

$$E(x, t) = E_{\text{body}}(x, t) \otimes E_{\text{neural}}(x, t)$$

The field is characterised by:

- **Multiplicity**: multiple emotional modes can be simultaneously active and interfering
- **Continuity**: it exists at all times, not only during episodes of conscious feeling
- **Spatial distribution**: different aspects of the field are localised in different regions of the soma (the familiar clinical observation that grief is felt in the chest, fear in the gut, anger in the jaw and fists)
- **Temporal dynamics**: the field evolves continuously, driven by the energy function

Figure 1. *The Soma-Field. The body and brain are not separate containers of emotion but two coupled components of a single distributed wave field. Neither is primary; each continuously modifies the other. The $\approx$ symbols indicate that wave activity is always present in each region, not only during episodes of conscious feeling.*

The Perception Threshold

Not all activity in the emotional field is consciously perceived. The field has a **perception threshold** T_i for each emotional mode i . Below this threshold, the emotional mode is sub-perceptual: it exists, it influences behaviour and physiology, but it does not surface as a named conscious feeling.

$$\text{Emotion } i \text{ is consciously perceived} \iff |\mathbf{E}_i(t)| > T_i$$

This threshold crossing corresponds precisely to the QFT excitation analogy: the emotional mode behaves like a virtual particle that has accumulated enough energy to become real — to emerge from the sub-threshold background and enter awareness.

This accounts for a range of clinically significant phenomena:

Clinical Observation	Soma-Field Account
Patient reports no feeling but shows physiological signs of distress	Sub-threshold field activity below T_i
Sudden unexpected flood of emotion in session	Rapid threshold crossing after gradual accumulation
Emotion felt somatically but not named	Threshold crossed in $\mathbf{E}_{\text{body}}$, not yet in $\mathbf{E}_{\text{neural}}$
Alexithymia (difficulty identifying feelings)	Elevated T_i — high threshold requiring more energy to cross
Hypervigilance / emotional flooding	Lowered T_i — reduced threshold, field crosses to conscious easily

Table 1. *Clinical observations mapped onto the perception threshold model.*

Figure 2. *The perception threshold T_i for a single emotional mode. The field is active continuously (lower trace). Conscious experience arises only when amplitude exceeds T_i (upper trace). Everything below the line is still there — shaping body and behaviour before it can be named.*

Figure o. *Continuous soma-field activity (blue) with a single threshold-crossing event. The field is always active; conscious experience (shaded) arises only when amplitude exceeds the perception threshold θ (red dashed). Below the threshold: real, causally active, but not yet conscious.*

The Interaction of Emotional Modes

Multiple emotional modes are simultaneously active in the field at all times. They do not simply co-exist: they interact. The nature of these interactions is encoded in the **emotional coupling matrix** W , where W_{ij} represents the influence of emotional mode j on emotional mode i .

- If $W_{ij} > 0$: emotion j amplifies emotion i (e.g., fear can amplify shame)
- If $W_{ij} < 0$: emotion j suppresses emotion i (e.g., calm suppresses anxiety)
- If $W_{ij} = 0$: emotions i and j are independent

The field evolves according to the energy gradient:

$$\dot{\mathbf{e}}(t) = -\nabla H(\mathbf{e}) + \eta(t)$$

where $\eta(t)$ represents the continuous low-level fluctuations of the sub-perceptual field — the emotional equivalent of quantum vacuum noise. The field is always moving, always seeking lower energy, never at absolute rest.

The Three-Layer Architecture

The nervous system that implements the soma-field is not architecturally flat. Three hierarchically organised layers contribute to field dynamics, each corresponding to a distinct evolutionary substrate and a distinct role in the model. The clinical literature (Porges, 2011; van der Kolk, 2014; Ogden et al., 2006) converges on this stratification; what follows is its formal expression.

Layer 1 — Brainstem / autonomic baseline. The oldest structures: vagal nuclei, arousal systems, interoceptive machinery. In the model, this layer is represented by the noise term and, specifically, by the heart rate variability coherence C_{HRV} , which modulates effective noise amplitude across the whole field:

$$\sigma_{\text{eff}} = \frac{\sigma_0}{C_{\text{HRV}}}$$

High HRV coherence narrows effective noise, stabilising the field in its current attractor. This is the mechanism of HRV biofeedback as a regulatory intervention: it does not target any specific emotional mode but lowers the fluctuation floor of the entire field.

Layer 1 extension: cardiac acceleration and landscape tilt. The term C_{HRV} measures the *current state* of cardiac regularity — where the heart is. A complementary quantity is $\dot{H}(t)$, the first time-derivative of heart rate, in units of beats/s². This is the **cardiac acceleration**: not what the heart rate is, but where it is going.

The dimensional parallel with gravity is exact: gravitational acceleration g carries units m/s²; cardiac acceleration $\dot{H}$ carries units beats/s². Both are accelerations; both describe a force field rather than a

position. Gravity does not tell you where a test mass is — it tells you how it will move next. Cardiac acceleration tells you not the current BPM but the direction of the next one: the $N+1$ state.

In the soma-field, $\dot{H}(t)$ enters the dynamics not as noise modulation but as a **landscape tilt** — a time-varying bias added to the Hamiltonian that tips the energy function toward activation or rest attractors:

$$H(\mathbf{e}, t) = H_0(\mathbf{e}) - \alpha \dot{H}(t) \beta \cdot \mathbf{e}$$

where $\alpha > 0$ is the cardiac-somatic coupling constant and β is a mode-coupling vector (at leading order, $\beta = \mathbf{1}$: the tilt acts uniformly across all modes). When $\dot{H}(t) > 0$ (heart accelerating), the landscape tilts toward higher activation states before any cognitive or affective threshold is crossed. When $\dot{H}(t) < 0$ (heart decelerating), it tilts toward rest. The full three-layer equation including the cardiac acceleration term is:

$$\dot{\mathbf{e}}(t) = -\nabla H_0(\mathbf{e}) + \alpha \dot{H}(t) \beta + R_{\text{PFC}}(\mathbf{e}, t) + \frac{\sigma_0}{C_{\text{HRV}}} \xi(t)$$

The two cardiac terms serve distinct functions: C_{HRV} (state) modulates the noise floor; $\dot{H}$ (acceleration) tilts the deterministic landscape. Both are needed for a complete account of cardiac influence on the field.

Predictive clinical value. A patient with $\text{BPM} = 90$ and $\dot{H} = +4 \text{ beats/s}^2$ is approaching threshold; one with $\text{BPM} = 90$ and $\dot{H} = -4 \text{ beats/s}^2$ is retreating from it. The snapshot is identical; the trajectories are opposite. Cardiac acceleration is therefore an early-warning signal for threshold crossings — detectable at Layer 1 before the emotional field at Layer 2 has crossed its threshold. This has independent support in cardiology: Bauer et al. (2006) demonstrated that *acceleration capacity* and *deceleration capacity* of heart rate — estimates of $\dot{H}$ over a cardiac window — carry prognostic information independent of conventional HRV measures.

The somatic equivalence principle. The cardiac acceleration term $\alpha \dot{H} \beta$ is structurally identical in the equation to any other forcing term. From the perspective of the field itself — from conscious experience — cardiac-driven activation is indistinguishable from event-driven activation. A sudden heart rate acceleration tilts the landscape by exactly the same mechanism as an external threat or an intrusive memory. The field has no access to the origin of the tilt. This is the formal account of a clinically well-documented phenomenon: anxiety initiated by cardiac irregularity (arrhythmia, postural hypotension, caffeine, exertion) is experienced as emotionally caused, because the somatic signal is identical. Disambiguation requires either external measurement or deliberate interoceptive inquiry that can distinguish the two sources.

Layer 2 — Limbic system / emotional memory. The primary substrate of the Soma-Field Model. The coupling matrix W , memory kernel $K(\tau)$, Hamiltonian $H(\mathbf{e})$, and threshold T all belong here. The limbic layer stores emotional-somatic states and reinstates them in response to partial body cues: a continuous, asymmetric, temporally extended Hopfield network operating on somatic states rather than cognitive patterns. This is the architectural layer that has been absent from every artificial neural network since McCulloch and Pitts (1943) (McCulloch & Pitts, 1943). The cortex has been modelled many times; the limbic system has not.

Structural plasticity under adversity. The Soma-Field framework permits a formal characterisation of the field's resilience under adverse conditions. Define the *plasticity index* Π as a composite of three measurable field properties:

$$\Pi = \frac{1}{S_{\text{inst}}} + \left. \frac{\partial \|W\|}{\partial t} \right|_{\text{adversity}} + C_{\text{HRV}}^{\text{recovery}}$$

The three terms correspond to: (i) how accessible regulated-state attractors remain under adversity ($1/S_{\text{inst}}$, instanton accessibility — Section 4.4); (ii) how much the coupling matrix can structurally adapt following a threshold crossing ($\partial \|W\|/\partial t$, the plasticity component); and (iii) how quickly the HRV floor recovers after activation ($C_{\text{HRV}}^{\text{recovery}}$, the regulatory resilience component). Complex PTSD is the clinical presentation of chronically low Π across all three terms simultaneously: high barriers to regulated attractors, a rigid W dominated by threat configurations, and impaired C_{HRV} recovery. Structural plasticity is the capacity of the field to update W in the aftermath of adversity without the adversity permanently *becoming* W .

Layer 3 — Neocortex / prefrontal regulatory layer. Top-down modulation of Layer 2, represented as a regulatory term $R_{\text{PFC}}(\mathbf{e}, t)$. The full field dynamics becomes:

$$\dot{\mathbf{e}}(t) = -\nabla H(\mathbf{e}(t)) + R_{\text{PFC}}(\mathbf{e}, t) + \frac{\sigma_0}{C_{\text{HRV}}} \xi(t)$$

R_{PFC} represents voluntary attention, therapeutic technique, and conscious reappraisal acting on the field. It is not a correction of Layer 2 but a modulation of it. Under sustained therapeutic engagement, R_{PFC} participates in the structural modification $W \rightarrow W'$ constituting the forward transformation (Section 7).

The **threshold T is the Layer 2 / Layer 3 boundary**: sub-threshold dynamics are processed limbically and remain below conscious awareness; threshold-crossing events enter Layer 3 and become available for narrative, meaning-making, and voluntary response. This is the formal basis for the clinical observation that insight without somatic activation is limited, and somatic activation without Layer 3 engagement cannot produce structural change: the layers are coupled, not independent. R_{PFC} requires a threshold crossing in order to have something to work with.

The two-term Langevin equation introduced in Section 3.3 is the Layer 2 special case ($R_{\text{PFC}} = 0$, $C_{\text{HRV}} = 1$). All subsequent sections develop that special case. The full three-layer equation is the general form.

The Energy Landscape

The Structure of the Emotional Energy Function

The energy function $H(\mathbf{e})$ defines a landscape over the space of possible emotional states. Like a physical landscape of hills and valleys, this landscape has:

- **Valleys (local minima)**: stable emotional states the field naturally moves toward
- **Hills (local maxima)**: unstable configurations the field naturally moves away from
- **Saddle points**: transitional configurations with mixed stability


The key property of an energy function is directionality: the field always moves *downhill*. It always evolves toward lower energy. Therapeutic intervention, in this framework, can be understood as:

1. **Changing the landscape**: modifying W — the coupling matrix — through new relational experience, insight, or somatic work, so that the energy minima are in healthier locations
2. **Adding energy to escape a trap**: helping the field accumulate enough energy to escape a deep but unhealthy local minimum (e.g., the freeze state)
3. **Pointing toward the global minimum**: orienting the field toward regulated calm

Attractor States: Fight, Flight, Freeze, and Regulated Calm

The Soma-Field Model proposes that the major attractor basins of the emotional energy landscape correspond directly to the autonomic states described by Porges' polyvagal theory.

Figure 3a. Topographic (bird's-eye) view of the energy landscape. The field always rolls downhill toward the nearest minimum. Freeze and calm are both low-energy — but freeze is surrounded by high walls. Escape from freeze requires crossing those walls, which means first gaining energy before losing it again. This is the clinical challenge of working with dissociative states.






Figure 3b. Schematic energy landscape. Fight/flight are high-energy, unstable local minima. Freeze is a low-energy but isolated attractor — easy to enter, hard to escape. Regulated calm is the global energy minimum.

Attractor	Energy State	Polyvagal Correlate	Clinical Presentation
Regulated Calm	Global minimum	Ventral vagal (social engagement)	Present, flexible, connected
Fight	Shallow high-energy minimum	Sympathetic (mobilisation)	Agitation, anger, urgency
Flight	Saddle point / shallow minimum	Sympathetic (mobilisation)	Anxiety, avoidance, rumination
Freeze	Deep isolated minimum	Dorsal vagal (immobilisation)	Dissociation, numbness, collapse

Table 2. Emotional attractors mapped onto Polyvagal states.

The therapeutic significance of this structure is considerable. The freeze state is dangerous not because it is high-energy — it is in fact very low energy — but because it is *isolated*: surrounded by energy barriers that make it difficult to exit. Escape from freeze requires first *increasing* the field's energy (mobilising some arousal) before it can flow toward regulated calm. This corresponds well to the clinical observation that working with dissociated patients requires careful titration of arousal — not too much, not too little — before emotional processing is possible.

The Coupling Matrix as a Personal Signature

The coupling matrix W is not universal. Each person has a unique W , shaped by attachment history, trauma, cultural context, and temperament. A person with a history of developmental trauma may have a W in which anxiety and shame are strongly coupled ($W_{\text{shame, anxiety}} \gg 0$), creating a combined attractor that is particularly deep and sticky. A person with a secure attachment history may have a W in which positive emotions are broadly coupled to one another, creating a wide basin around regulated calm.

This implies that the energy landscape is a therapeutic object in its own right: understanding a patient's W is understanding the structural dynamics of their emotional life.

In the M-theory compactification analogy developed in Appendix A, the coupling topology W corresponds to the shape of the compact G_2 manifold — the seven-dimensional geometry that determines which force-like couplings are allowed and with what strengths. That analogy is here made precise: two people differ not merely in their emotional *parameter settings* but in their coupling *geometry*. Developmental trauma does not set a dial to the wrong value; it deforms the manifold. The therapeutic process of modifying W through relational experience, insight, or somatic work is, in this language, differential geometry: a continuous deformation of the G_2 manifold toward a configuration in which the regulated-calm attractor is globally accessible. The practitioner is, without having been told so, a geometer.

Dissonance and Resolution

The Acoustic Analogy

The Soma-Field Model draws a further structural analogy, this time with acoustics. When two sound waves interact, the quality of their interaction — consonance or dissonance — depends on the phase relationship between them. Consonant intervals (the octave, the fifth) have simple frequency ratios and produce stable, reinforcing interference patterns. Dissonant intervals (the tritone, the minor second) have complex ratios and produce beating, unstable, tension-generating patterns.

The model proposes that the same relationship holds between emotional modes. When two emotional modes are in a compatible relationship — when their interaction is consonant — the field is in a relatively low-energy configuration and moves naturally toward the energy minimum. When they are in an incompatible relationship — when their interaction is dissonant — the field is in a higher-energy configuration, generating a gradient that drives toward resolution.

Dissonance, in this framework, is felt as tension. It is not pathological; it is directional. Dissonance is the field's way of communicating that it is far from equilibrium and that resolution is available.

The Resolution Principle

In music, dissonance resolves to consonance. The tritone — the most dissonant interval in Western tonality — creates a powerful gravitational pull toward resolution. In counterpoint, the rules of voice leading describe the specific paths by which dissonance must resolve. These rules are not arbitrary conventions; they describe the geometry of the acoustic energy landscape.

The same principle applies to emotional dissonance. An unresolved emotional state — grief that has not been fully experienced, anger that has been suppressed, fear that has been dissociated — is a dissonance in the field. It generates a persistent tension gradient. The therapeutic process can be understood as guided voice leading: finding the specific path of resolution that transforms the dissonant configuration into a consonant one.

This provides a formal basis for a widely-held clinical intuition: that emotions need to be *felt through* rather than avoided. Avoidance keeps the field in a dissonant state. The energy minimum — regulated calm — lies on the other side of the dissonance, not around it.

The Soma-Field Instrument

Rationale

The Soma-Field Model is not only a theoretical framework. It motivates a practical therapeutic instrument: a means by which a person can *externalise* their emotional field — make it visible and audible — and interact with it in real time.


The core insight is that the emotional field is normally invisible to its host. It operates below the threshold of conscious awareness, shaping behaviour and physiology without being available for reflection. If its activity could be rendered as a signal — a sound, an image, a pattern — it could become an object of therapeutic attention.

Design

The instrument uses a MIDI controller with 16 rotary knobs as its input interface. Eight emotional dimensions are encoded, each represented by two knobs:

- **Knob 1** of each pair: the somatic (body-level) intensity of that emotional mode
- **Knob 2** of each pair: the cognitive/neural intensity of that emotional mode

This design reflects the two-component structure of the field: body and mind are encoded separately but coupled in the computation. Each knob has a continuous range, allowing fine expression of emotional intensity.






Figure 4. The Soma-Field Instrument: input, computation, and multimodal output.

The Feedback Loop

The instrument creates a **closed feedback loop** between the person and their emotional field:

1. The person expresses their current emotional state by adjusting the knobs
2. The system computes the energy function $H(\mathbf{e})$ and its gradient ∇H
3. The energy level, dissonance, and proximity to attractor states are rendered as:
 - **Sound:** harmonic content and timbre reflect the consonance or dissonance of the current state
 - **MIDI output:** pitch rises with tension, resolves as energy decreases
 - **Visual:** a real-time map of the emotional field, showing wave activity, threshold crossings, and the direction of the energy gradient
4. The person hears and sees their emotional field, and adjusts the knobs in response

This loop externalises the emotional field's gradient — the direction in which it is *trying* to move — and makes it available as sensory information. The person becomes not only the source of the emotional signal but also its observer, creating the conditions for reflection and regulation that are at the heart of therapeutic work.

The Pluggable Emotion Model

No single model of the emotions is assumed. The coupling matrix W — the structure that determines how emotional modes interact — is loaded from an external configuration file. Standard models (Plutchik's wheel of emotions, Ekman's basic emotions, the valence-arousal-dominance dimensional model) are provided as defaults. The therapist or client can modify the coupling values to reflect their own understanding of their emotional patterns, or a new model can be substituted entirely. The computational engine is model-agnostic.

Clinical Implications

Assessment

The Soma-Field Model suggests a different orientation for emotional assessment. Rather than asking “What emotion do you feel?” — which presupposes threshold-level conscious awareness — it invites attention to the sub-perceptual field: “What is present in the body right now, even if you cannot name it?” This aligns with Focussing-oriented approaches and with sensorimotor methods that prioritise somatic signal over narrative content.

The energy landscape provides a clinical map. A person chronically in a fight or flight attractor shows a different energy signature from a person in a freeze attractor, even if their presenting narratives are superficially similar. The model suggests that these are structurally different therapeutic challenges: fight/flight require down-regulation, while freeze may first require careful up-regulation before down-regulation becomes possible.

Intervention

The energy function provides a formal basis for several existing clinical interventions:

- **Grounding and titration** (Levine, 2010): adding small, controlled amounts of energy to the field to approach — without flooding — a previously frozen or avoided emotional state
- **Pendulation** (Levine, 2010): oscillating between a dissonant state and a resource state, progressively widening the tolerance window — equivalent to approaching the energy minimum via a series of small excursions
- **Somatic resourcing** (Ogden et al., 2006): establishing a stable low-energy region in the landscape that the field can return to after excursions into high-energy territory
- **Working with the felt sense** (Gendlin, 1978): attending to sub-threshold field activity and allowing it to cross the perception threshold in a supported context

Psychoeducation

The wave model is immediately accessible to clients who have struggled to understand their emotional experience. The statement: “Your emotions are like waves — they are always there, even when you can’t feel them, and they are always moving” is both technically accurate within the Soma-Field framework and clinically useful as a normalising frame for sub-threshold emotional activity, for the apparently sudden onset of strong feelings, and for the experience of feeling multiple conflicting emotions simultaneously.

The energy landscape metaphor — “right now the field is in a valley that is hard to leave, but it is not the lowest valley available to you” — offers a way to discuss the freeze state, dissociation, and emotional

stuckness without pathologising, while still acknowledging the structural difficulty of these states and the work required to shift them.

Neurodivergent Conditions as Operator Modifications


A clinically significant extension of the Soma-Field Model concerns neurodivergent conditions — specifically Autism Spectrum Condition (ASC), Attention Deficit Hyperactivity Disorder (ADHD), and Complex Post-Traumatic Stress Disorder (C-PTSD), which frequently co-occur and each present distinct challenges for somatic emotional processing.

The key architectural principle is this: **these conditions are not parameter settings in the model. They are structural modifications to the operators themselves.** This distinction matters both mathematically and clinically. A parameter change (“set the fear threshold lower”) is a quantitative adjustment within the existing structure. An operator modification changes the *form* of the dynamics — it alters the governing equations, not merely their coefficients. Each condition wraps the standard pipeline in a different functional modifier, and — critically for the many individuals who carry all three — these modifiers *compose*. The combined condition is not three separate problems; it is the composition of three operators acting on the same underlying field.

The mathematical details of each modifier are given in Appendix B. Clinically, the consequences are as follows.


Complex PTSD introduces a *memory kernel* into the field dynamics: past high-energy states leave decaying echoes that continue to excite the field without new external stimulus. This is why traumatic activation can appear without identifiable trigger — the field is responding to its own history, not its current environment. The standard Hopfield attractor topology is also disrupted: C-PTSD renders the freeze attractor pathologically deep and wide, the window of tolerance (the basin around regulated calm) pathologically narrow, and the coupling matrix asymmetric — a condition under which the field can enter persistent *limit cycles* rather than settling to a stable minimum. Re-experiencing, flashbacks, and hypervigilance are, in this framework, limit-cycle oscillations in the traumatised field.

REGULATED FIELD (symmetric W , no memory kernel)



- ↑ baseline returns to near-zero between episodes
- ↑ each threshold crossing is a discrete, independent event
- ↑ 'regulated calm' is a genuine resting state – the global energy minimum

C-PTSD MODIFIED FIELD (asymmetric W , memory kernel $K(t-s)$ present)



- ↑ memory kernel: each activation feeds energy back into the next
- ↑ field rarely returns to true rest – past states re-enter present dynamics
- ↑ almost entirely above T : activation is the default, not the exception
- ↑ 'regulated calm' requires a non-perturbative transition (the instanton):
small steps do not reach it; a qualitatively different move is needed

Figure 5. The same emotional field mode under two dynamic regimes. Top: regulated dynamics — the field oscillates and returns to a low baseline between episodes; conscious emotion (above T) is episodic and resolves. Bottom: C-PTSD-modified dynamics — the memory kernel elevates the baseline so that the field rarely returns to rest; episodes bleed into one another; the system cycles rather than settles. The mathematical basis for this comparison is given in Appendix B.2.

Developmental timing and what can be recovered. The character of the C-PTSD modification depends critically on *when* it occurred — the developmental age τ_d at which the primary traumatic modification took place.

For **late trauma** (τ_d large — adult or post-verbal): a coupling matrix W_0 formed before the event. The modification is additive: $W = W_0 + \delta W_{\text{trauma}}$. A counterfactual pre-trauma self exists, encoded in explicit narrative memory. Therapeutic processing can target δW specifically, and the goal of recovering proximity to W_0 is formally coherent.

For **early trauma** (τ_d small — pre-verbal, perinatal): the coupling matrix W was *formed under the modification*. There is no W_0 . The asymmetric coupling and the memory kernel coefficients are the baseline architecture, not additions to one. A counterfactual pre-trauma self was never encoded — it does not exist as a recoverable state.

This is a formal statement of a clinical fact that somatic therapists recognise but rarely have a mechanistic basis for: early trauma cannot be *processed away* in the sense of recovering a prior self, because no

prior self was formed. The therapeutic goal is not subtraction ($W \rightarrow W_0$, which is undefined) but **forward transformation**: constructing a W' that supports a wider window of tolerance, different attractor topology, and lower memory-kernel amplitudes. This is a different mathematical operation — and requires a different therapeutic model.

The Soma-Field Instrument can reflect this distinction directly: a user whose primary modification is pre-verbal initialises with a *structural* coupling matrix (the modification *is* the baseline), not a neurotypical matrix with an added modifier. The formal basis for this parameterisation is given in Appendix B.2.1.

ADHD raises the effective *thermal noise* of the field — the amplitude of the sub-perceptual fluctuations — and simultaneously reduces the damping coefficient that slows the field's response to the energy gradient. The result is a field that explores its energy landscape rapidly and unpredictably, is easily displaced from shallow attractor basins by small perturbations (distractibility), but also achieves states of intense concentration (hyperfocus) when the coupling to a high-salience stimulus temporarily deepens a specific attractor basin far beyond its resting depth. ADHD is not a deficit of attention; it is a high-temperature, low-damping emotional field with a stimulus-dependent attractor structure.

Autism Spectrum Condition modifies the *projection kernels* — the functions that determine how the continuous somatic field is sampled to produce the discrete state vector — and the *sparsity* of the coupling matrix. Interoceptive research in autism (Garfinkel et al., 2016) documents significant differences in the processing of internal body signals; in model terms, certain somatic regions are over-represented (heightened sensory sensitivity) and others under-represented (reduced interoceptive clarity, contributing to alexithymia). The coupling matrix in ASC tends toward greater sparsity — fewer strong cross-modal emotional couplings — a pattern consistent with monotropism (Murray, 2018): the field settles deeply into individual attractors but transitions between them require proportionally more energy. Intense interests, emotional consistency within a context, and difficulty with unexpected transitions all follow from this attractor topology.

For the Soma-Field Instrument, the practical implication is significant. Rather than asking a neurodivergent user to configure their experience through knob adjustments, the system can instantiate the appropriate operator modifications as a named profile — “load C-PTSD modifier”, “load ADHD modifier” — each of which transforms the pipeline at the correct mathematical level. The user then interacts with a field that already reflects their structural reality, rather than one calibrated for a neurotypical baseline.

A further clinical implication deserves explicit statement. The Soma-Field Model locates interoceptive accuracy in the field itself: whether a somatic signal has exceeded its perceptual threshold T_i is a property of the field state, not a property of the clinician's assessment of the patient's credibility. A patient reporting an acute somatic state is reporting a threshold-crossing event. The model provides

no mechanism by which external disbelief suppresses that crossing. Modified projection operators — as occur in ASC — produce *different* somatic self-reports; the model gives no reason to assume they produce *less accurate* ones. The clinical literature documents a systematic tendency to interpret unusual interoceptive self-reports from neurodivergent patients as indicative of psychogenic origin rather than genuine somatic signal (Nicolaidis et al., 2015). The Soma-Field Model predicts that this interpretive pattern constitutes a category error: it confuses operator modification with signal absence. The practical consequences — missed diagnoses, deferred treatment, and the iatrogenic re-inforcement of existing trauma — are well-documented and, within this framework, mathematically predictable.

Limitations and Future Directions

The Soma-Field Model is a theoretical framework and must be evaluated as such. Its current form makes several idealisations that require scrutiny.

The coupling matrix W is treated as a fixed parameter, but emotional coupling is dynamic: it changes with context, relationship, and developmental history. A more complete model would treat W as a slowly-evolving quantity, shaped by the field's own history — a form of synaptic plasticity applied to the emotional domain.

The threshold T_i is treated as a fixed property of each emotional mode, but experimental evidence suggests that thresholds are modulated by attentional focus, arousal level, and interpersonal context. A person in a safe therapeutic relationship will typically have lower thresholds — more material reaches conscious awareness — than the same person in an unsafe context.

The acoustic analogy, while structurally productive, requires empirical grounding. The claim that emotional dissonance and acoustic dissonance share formal properties is a hypothesis, not an established finding. Empirical work comparing physiological measures of emotional tension with acoustic analysis of synchronised vocal or musical output would be a productive direction for testing this claim.

The instrument described in Section 6 is a prototype concept. User studies with clinical populations, and collaboration with practising therapists, will be required to assess its therapeutic utility and to identify appropriate clinical contexts.

Future theoretical work should address the relational field: the observation, familiar in systemic and relational approaches to psychotherapy, that emotional fields are not bounded by individual bodies but are co-generated in the space between people. The coupling matrix W of a relationship may be as clinically significant as the W of an individual.

Axiomatic QFT status (update, 2026). A subsequent paper in this series (P14, *The Universal Somatic Field as a Euclidean Quantum Field Theory*) proves that the free-field USF satisfies all five Osterwalder–Schrader axioms, placing it within the rigorous framework of constructive quantum field theory. The proof is machine-verified in Lean 4 with zero sorries. Reflection positivity (OS₃) guarantees the legitimacy of the Minkowski continuation proved in the temporal-dynamics companion paper. The interacting (Hopfield-coupled) theory is addressed in P15.

Conclusion

The Soma-Field Model proposes a formally grounded account of emotional dynamics that is consistent with the clinical observations of somatic psychotherapy, polyvagal theory, and Focussing-oriented practice. Its central claims — that emotions are a persistent distributed field, that conscious experience is a threshold crossing, and that emotional dynamics are governed by an energy function that drives the field toward stable attractor states — are not novel as clinical intuitions. What is novel is the formal structure that unifies them, and the instrument that the structure motivates.

The model does not resolve the philosophical question of what emotions fundamentally *are*. It offers instead a working representation: one that is precise enough to be computationally implemented, close enough to existing clinical frameworks to be therapeutically applicable, and open enough to be modified as understanding deepens. It invites the therapist to think of the consulting room as a space in which two emotional fields interact — each shaping the other’s energy landscape — and of therapeutic work as the art of attending to that interaction with enough precision and care to guide both fields toward lower energy, toward greater coherence, toward regulated calm.

The wave is always there. Therapy is learning to listen to it.

A note on provenance. The Soma-Field Model was not developed from a position of theoretical neutrality. The author carries, as primary data, a lifetime of direct experience of the dynamics described above. The neurodivergent operator modifications of Appendix B are not theoretical abstractions: the C-PTSD memory kernel of B.2 was installed pre-verbally, at approximately eighteen months of age, during a developmental trauma that predates language acquisition entirely. No narrative trace of the origin event exists — there was no verbal capacity with which to encode one. Only the field echo remains, and a measurable physical asymmetry in the body that received it. The ASD and ADHD operator modifications of Appendix B.4 and B.3, respectively, were the instruments by which the model was subsequently constructed: the monotropic attractor structure of B.4 provided the capacity for sustained engagement with an entirely unfamiliar theoretical domain; the high-temperature field dynamics of B.3 drove rapid traversal across it.

The proximate cause is described in full in the companion patient-facing publication. Briefly: an acute somatic emergency in 2025 — a genuine threshold-crossing event, later confirmed as cerebral hypoxia secondary to Long Covid — was attributed, at clinical presentation, to psychiatric origin. The present paper is, among its other functions, a formal response to that attribution.

The causal chain is as follows. A pre-verbal trauma in approximately 1968 installed the C-PTSD operator modifications described in Appendix B.2. The ASD and ADHD modifications of Appendix B.3 and B.4 shaped the system across the intervening decades. Fifty-seven years later, that system's accurate interoceptive signal was dismissed as psychiatric noise. The paper which formally demonstrates that this dismissal constitutes a category error was produced, as a direct causal consequence, by the same operator stack that it describes. The paper is the fixed point of its own subject matter. The author considers this observation methodologically significant.

Publication Claim Registry

To support claim-level review rather than all-or-nothing acceptance, this manuscript registers its highest-impact claims with scope labels and disconfirmation tests.

Claim ID	Claim	Scope	Evidence in this work	Disconfirmation criterion
SF-1	Conscious percept is a propagator pole of the soma-field	S1 Structural	Formal derivation in Sections 2-3	Inability to express percept dynamics as Green's-function response under stated operator
SF-2	Emotional attractors are Hopfield-energy minima	S2 Predictive	Energy model and trajectory framework	Constructed update rule under model assumptions with systematic energy ascent

Claim ID	Claim	Scope	Evidence in this work	Disconfirmation criterion
SF-3	Threshold governs felt vs sub-felt emotional activity	S2 Predictive	Threshold operator and clinical mapping	Reliable high-amplitude mode activity with no threshold-dependent behavioural or physiological signature
SF-4	Topological barriers explain classical therapeutic plateaus	S2 Predictive	Formal treatment plus linked companion experiments	Controlled demonstration that matched low-noise classical dynamics crosses registered barriers at equivalent rate
SF-5	Quantum extension yields topological reachability advantage	S2 Predictive	QUANT-EXP-1 companion evidence and linked artifacts	Controlled replication showing no reachability advantage over matched classical baseline

Scope labels: S₁ = structural; S₂ = predictive; S₃ = independently replicated. Current publication target for core claims is S₂.

Claim-Evidence-Result Matrix

To make review traceable, each core claim is paired with concrete evidence outputs and current result status.

Claim ID	Evidence artifact(s)	Current result status
SF-1	Sections 2-3 derivation of field/propagator structure	structural derivation complete
SF-2	Energy formulation + instrument runtime equations	predictive structure complete
SF-3	Threshold operator definition + clinical interpretation sections	predictive mapping complete
SF-4	Barrier analysis; companion paper <i>Quantum Soma and the Penrose Gap</i> (doi:10.5281/zenodo.20351230)	confirmed (QUANT-EXP-1 PASS)
SF-5	QUANT-EXP-1 experiment outputs (see supplementary archive, doi:10.5281/zenodo.20351230)	confirmed: cold 0/200, CI [0.000, 0.019]; quantum peak 0.408–0.410; all hardening checks PASS

This matrix is intended for reviewer navigation and is updated as companion results are expanded or independently replicated.

Replication Package Requirements

To make SF-2 through SF-5 externally testable, each release tagged for review must ship a minimal replication package that can be executed without private context.

Required contents:

1. simulation code and support modules (see supplementary archive, doi:10.5281/zenodo.20351230),
2. full parameter snapshot ($W, \mathbf{b}, \gamma, D, \theta$, temperature policy),
3. raw trajectory logs with timestamped attractor labels,
4. analysis scripts that produce the reported summary tables,
5. frozen output artifacts (CSV/plots) referenced in this manuscript.

A claim remains S2 until an independent operator reproduces directionally consistent outcomes from this package under the same declared protocol.

Reviewer-Risk Objections and Responses

To reduce ambiguity in peer review, the highest-probability objections are mapped to bounded responses and concrete upgrade paths.

Reviewer objection	Current response in this manuscript	Remaining action to reach stronger status
“This is an analogy, not a formal model.”	Sections 2-4 define operators, dynamics, and testable predictions; Section 9.1 registers disconfirmation criteria claim-wise.	Promote more claims from S2 to S3 via independent replication.
“Evidence is pilot-stage and may not generalize.”	Section 9.2 explicitly labels pilot support and companion-only scope.	Add multi-operator replication and blinded protocol variants.
“Quantum advantage may be implementation-specific.”	SF-5 includes a controlled disconfirmation criterion against matched classical baselines.	Publish full benchmark harness with pre-registered acceptance thresholds.
“Clinical interpretation may exceed data scope.”	Scope labels (S1/S2/S3) and claim registry separate structural from predictive claims.	Add prospective cohort evidence before any clinical-effectiveness claim.

Independent Replication Ledger Linkage

S2 to S3 promotion for this manuscript is governed by an independent replication ledger maintained in the supplementary archive (doi:10.5281/zenodo.20350515).

Tracked claim IDs in ledger scope: SF-2, SF-3, SF-4, SF-5.

Promotion gate: a claim is upgraded only when at least one ledger row records an independent operator PASS with a reproducible package hash and linked raw/derived evidence artifacts.

Acknowledgements

This work exists because ten years of psychotherapy moved the barriers far enough that two events in early 2026 could cross them. The theory is, among other things, a record of that.

Conclusion: The Integrated Account

The Oxford PPE tradition was built on the belief that philosophy, politics, and economics illuminate each other — that a student who has engaged seriously with all three is better equipped to understand the world than a student who has specialised in one. This book has made the stronger claim: that the three disciplines are not merely illuminated by each other but governed by the same mathematical structure, and that structure is the Universal Somatic Field.

The conclusion to this volume is not a summary but an integration: a sketch of what a field-theoretic PPE would look like as a mature discipline.

The Common Structure

Each of the three disciplines, in the field-theoretic account, reduces to a version of the same question: what is the attractor landscape, and what dynamics govern the approach to equilibrium?

Philosophy asks: what is the good for a person, and what dynamics lead a person toward or away from it? The field-theoretic answer: the good is the configuration of the somatic field in which the attractor landscape is healthy — rich in adaptive basins, free from trauma wells, ergodic and temperature-appropriate. The dynamics of the good life are the dynamics of the somatic field.

Politics asks: what is the just organisation of collective life, and what dynamics lead a society toward or away from it? The field-theoretic answer: just organisation is the institutional landscape in which the constraint field is ergodic (rule of law), rights are topologically stable (topological invariants), and the regulatory interventions are above the WKB threshold (effective). The dynamics of just governance are the dynamics of the social somatic field.

Economics asks: what is the efficient organisation of production and exchange, and what dynamics lead a market toward or away from it? The field-theoretic answer: efficiency is attractor optimality in the Hopfield network — the condition that the energy landscape has its minimum in the welfare-optimal configuration. The dynamics of market efficiency are the Hopfield dynamics.

The three questions have the same mathematical form. The answers are the same mathematical object: an attractor landscape, and the field dynamics that determine which attractors are reached and how quickly.

The Patient as Political-Economic Subject

The patient perspective paper in this volume served as the humanising chapter — the reminder that the abstract machinery of field theory has consequences for individual human lives. The patient is simultaneously a philosophical subject (a person with a somatic field, a phenomenal life, a good or bad existence), a political subject (a citizen whose rights are or are not protected, whose legal landscape is or is not ergodic), and an economic subject (a participant in markets whose structure may or may not be conducive to human flourishing).

The field-theoretic framework describes all three simultaneously. The same equation that describes the patient's trauma well describes the legal fragmentation that makes their rights insecure, and the market failure that makes their economic position precarious. This is not a coincidence. It is the recognition that the patient is embedded in field dynamics at multiple scales — personal, social, political, economic — and that all of these scales are governed by the same master equation.

What a Field-Theoretic PPE Would Teach

A PPE curriculum built around the field-theoretic framework would teach students to:

Identify the energy landscape. Before analysing any political, economic, or philosophical problem, identify the attractor landscape: where are the stable states, where are the barriers, what is the temperature of the field? This is the first move in every field-theoretic analysis.

Read the dynamics. Understand what the field dynamics predict: where will the system end up if left alone, what interventions are above threshold, what topological transitions are approaching? The dynamics answer the question that equilibrium theory cannot: what happens next?

Design for robustness. Evaluate any proposed institution, policy, or practice not just for efficiency in equilibrium but for topological robustness: will it persist under perturbation? Will it maintain the rule of law (ergodicity) under stress? Will it protect rights (topological invariants) against powerful countervailing forces?

Ground abstraction in lived experience. The patient perspective is not a soft addition to the hard analysis. It is the check on the analysis: does the abstract account describe a human life as it is actually lived? The somatic field dynamics are ultimately the dynamics of felt experience, and any account that loses sight of that has lost the point.

Open Questions for PPE

The field-theoretic PPE opens a set of questions that cross the traditional disciplinary boundaries.

The good life and field health. What would a positive account of field health look like — one that goes beyond the absence of trauma wells to specify what a flourishing somatic field configuration is? This is a philosophical question with empirical content: it requires both a normative theory and a measurement programme.

Institutional design for robustness. Given the field-theoretic account of institutional stability, what design principles produce institutions that are robustly ergodic and rights-preserving? This is a political theory question with engineering content.

Market design for landscape optimality. Given the Hopfield account of market dynamics, what market structures produce energy landscapes whose global minimum corresponds to welfare-optimal allocations? This is an economics question with physics content.

Mind, market, and mandate. The field is the common ground.

Berger, M. (1955). Sur les groupes d'holonomie homogène des variétés à connexion affine et des variétés riemanniennes. *Bulletin de La Société Mathématique de France*, 83, 279–330.

Black, F., & Scholes, M. (1973). The pricing of options and corporate liabilities. *Journal of Political Economy*, 81(3), 637–654. <https://doi.org/10.1086/260062>

Hertz, J., Krogh, A., & Palmer, R. G. (1991). *Introduction to the theory of neural computation*. Addison-Wesley.

Hopfield, J. J. (1982). Neural networks and physical systems with emergent collective computational abilities. *Proceedings of the National Academy of Sciences*, 79(8), 2554–2558. <https://doi.org/10.1073/pnas.79.8.2554>

Jaynes, E. T. (1957). Information theory and statistical mechanics. *Physical Review*, 106(4), 620–630. <https://doi.org/10.1103/PhysRev.106.620>

Johnson, A. (2026). *The Soma-Field: A wave-based model of emotional dynamics and its clinical implications*. <https://doi.org/10.31234/osf.io/h74e9>

Joyce, D. (2000). *Compact manifolds with special holonomy*. Oxford University Press.

Juslin, P. N., & Sloboda, J. A. (Eds). (2011). *Handbook of music and emotion: Theory, research, applications*. Oxford University Press.

McCulloch, W. S., & Pitts, W. (1943). A logical calculus of the ideas immanent in nervous activity. *The Bulletin of Mathematical Biophysics*, 5(4), 115–133. <https://doi.org/10.1007/BF02478259>

Moura, L. de, & Ullrich, S. (2021). The Lean 4 theorem prover and programming language. In A. Platzer & G. Sutcliffe (Eds), *Automated deduction – CADE 28* (Vol. 12699, pp. 625–635). Springer. https://doi.org/10.1007/978-3-030-79876-5_37

Nakahara, M. (2003). *Geometry, topology and physics* (2nd edn). Institute of Physics Publishing.

nLab authors. (2024). *nLab*. <https://ncatlab.org/nlab/show/HomePage>

Penrose, R. (1971). Applications of negative dimensional tensors. In D. J. A. Welsh (Ed.), *Combinatorial mathematics and its applications* (pp. 221–244). Academic Press.

Ramachandran, V. S., & Hirstein, W. (1998). The perception of phantom limbs: The D.O. Hebb lecture. *Brain*, 121(9), 1603–1630. <https://doi.org/10.1093/brain/121.9.1603>

Randall, L., & Sundrum, R. (1999). A large mass hierarchy from a small extra dimension. *Physical Review Letters*, 83(17), 3370–3373. <https://doi.org/10.1103/PhysRevLett.83.3370>

Selinger, P. (2010). A survey of graphical languages for monoidal categories. In B. Coecke (Ed.),

- New structures for physics* (Vol. 813, pp. 289–355). Springer. https://doi.org/10.1007/978-3-642-12821-9_4
- Seth, A. K. (2021). *Being you: A new science of consciousness*. Dutton.
- Snow, C. P. (1959). *The two cultures and the scientific revolution*. Cambridge University Press.
- The Univalent Foundations Program. (2013). *Homotopy type theory: Univalent foundations of mathematics*. Institute for Advanced Study. <https://homotopytypetheory.org/book/>
- Veneziano, G. (1968). Construction of a crossing-symmetric, Regge-behaved amplitude for linearly rising trajectories. *Il Nuovo Cimento A*, 57(1), 190–197. <https://doi.org/10.1007/BF02824451>
- Wilson, K. G., & Kogut, J. (1974). The renormalization group and the ϵ expansion. *Physics Reports*, 12(2), 75–199. [https://doi.org/10.1016/0370-1573\(74\)90023-4](https://doi.org/10.1016/0370-1573(74)90023-4)
- Zinn-Justin, J. (2002). *Quantum field theory and critical phenomena* (4th edn). Oxford University Press.